

Geethanjali College of Engineering and Technology
(UGC AUTONOMOUS INSTITUTION)

(Accredited by NBA and NAAC with 'A' grade, Approved by AICTE New Delhi and Affiliated to JNTUH)

Cheeryal (V), Keesara (M), Medchal (Dist.), Telangana – 501 301.

DATA ANALYTICS LAB
(18CS41L1)

Laboratory Manual

IV Year B.Tech. CSE I Semester



Geethanjali

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

2021-2022

Lab Incharge

HOD-CSE

Dr. A. Sree Lakshmi

Geethanjali College of Engineering and Technology
(UGC AUTONOMOUS INSTITUTION)

(Accredited by NBA and NAAC with 'A' grade, Approved by AICTE New Delhi and Affiliated to JNTUH) Cheeryal (V), Keesara (M), Medchal (Dist.), Telangana – 501 301.



Geethanjali

DATA ANALYTICS LAB

(18CS41L1)

Laboratory Workbook

Name: _____

Roll No: _____ **Class:** _____ **Branch:** _____

Academic Year: 2021 – 2022

Geethanjali College of Engineering and Technology
(UGC AUTONOMOUS INSTITUTION)

**(Accredited by NBA and NAAC with 'A' grade, Approved by AICTE New
Delhi and Affiliated to JNTUH)**

Cheeryal (V), Keesara (M), Medchal (Dist.), Telangana – 501 301.



Geethanjali

CERTIFICATE

*This is to certify that Mr. / Miss _____ has satisfactorily
completed _____ number of experiments in the **Data Analytics Laboratory**.*

Roll No: _____
Section: _____
Academic Year: _____

Branch: _____
Year: _____

Head
Dept. of CSE

Faculty
In-Charge

Internal Examiner

External Examiner

GEETHANJALI COLLEGE OF ENGINEERING AND TECHNOLOGY

(Autonomous)

Cheeryal (V), Keesara (M), Medchal Dist., Telangana-501301

18CS41L1-DATA ANALYTICS LAB

IV Year. B.Tech. (CSE) – I Sem

| L | T | P/D | C |
|---|---|-----|---|
| - | - | 2/- | 1 |

Prerequisite(s):

- 18CS2102 - Object Oriented Programming using Java
- 18CS2203 -Database Management Systems

Course Objectives:

Develop ability to

1. Know the basic elements of Big Data and Data science to handle huge amount of data.
2. Gain knowledge of basic mathematics behind the Big data.
3. Understand the different Big data processing technologies.
4. Apply the Analytical concepts of Big data using R and Python.
5. Visualize the Big Data using different tools.

Course Outcomes (COs):

At the end of the course, student would be able to:

- CO1: Observe Big Data elements and Architectures.
- CO2: Apply different mathematical models for Big Data.
- CO3: Demonstrate their Big Data skills by developing different applications.
- CO4: Apply each learning model for different datasets.
- CO5: Analyze needs, challenges and techniques for big data visualization.

LIST OF EXPERIMENTS

Week 1: Installation, Configuration, and Running of Hadoop and HDFS.

Week 2: Implementation of Word Count / Frequency Programs using MapReduce.

Week 3: Implementation of MR Program that processes a Weather Dataset.

Week 4: Implementation of Linear and Logistic Regression.

Week 5: Implementation of SVM Classification Technique.

Week 6: Implementation of Decision Tree Classification Technique.

Week 7: Implementation of Hierarchical Clustering.

Week 8: Implementation of Partitioning Clustering.

Week 9: Data Visualization using Pie, Bar, Boxplot Chart Plotting Framework.

Week 10: Data Visualization using Histogram Plotting Framework.

Week 11: Data Visualization using Line Graph Plotting, Scatterplot Plotting Framework.

Week 12: Application to analyze Stock Market Data using R Language.

Vision of the Institution

Geethanjali visualizes dissemination of knowledge and skills to students, who eventually contribute to well-being of the people of the nation and global community.

Mission of the Institution

- To impart adequate fundamental knowledge in all basic science and engineering technical and Inter - personal skills so students
- To bring out creativity in students that would promote innovation, research and entrepreneurship.
- To Preserve and promote cultural heritage, humanistic and spiritual values promoting peace and harmony in society.

Vision of the Department

To produce globally competent and socially responsible computer science engineers contributing to the advancement of engineering and technology which involves creativity and innovation by providing excellent learning environment with world class facilities

Mission of the Department

- To be a center of excellence in instruction, innovation in research and scholarship, and service to the stake holders, the profession, and the public.
- To prepare graduates to enter a rapidly changing field as a competent computer science engineer.
- To prepare graduate capable in all phases of software development, possess a firm understanding of hardware technologies, have the strong mathematical background necessary for scientific computing, and be sufficiently well versed in general theory to allow growth within the discipline as it advances.
- To prepare graduates to assume leadership roles by possessing good communication skills, the ability to work effectively as team members, and an appreciation for their social and ethical responsibility in a global setting.

PROGRAM EDUCATIONAL OBJECTIVES

- To provide graduates with a good foundation in mathematics, sciences and engineering fundamentals required to solve engineering problems that will facilitate them to find employment in industry and / or to pursue postgraduate studies with an appreciation for lifelong learning.
- To provide graduates with analytical and problem-solving skills to design algorithms, other hardware / software systems, and inculcate professional ethics, inter-personal skills to work in a multi-cultural team.

- To facilitate graduates to get familiarized with the art software / hardware tools, imbibing creativity and innovation that would enable them to develop cutting-edge technologies of multi-disciplinary nature for societal development.

PROGRAM OUTCOMES (POs)

Program Outcomes (POs) describe what students are expected to know and be able to do by the time of graduation to accomplish Program Educational Objectives (PEOs). The Program Outcomes for Computer Science and Engineering graduates are:

Engineering Graduates would be able to:

PO 1: Engineering knowledge: Apply the knowledge of mathematics, science, engineering fundamentals, and an engineering specialization to the solution of complex engineering problems.

PO 2: Problem analysis: Identify, formulate, review research literature, and analyze complex engineering problems reaching substantiated conclusions using first principles of mathematics, natural sciences, and engineering sciences.

PO 3: Design/development of solutions: Design solutions for complex engineering problems and design system components or processes that meet the specified needs with appropriate consideration for the public health and safety, and the cultural, societal, and environmental considerations.

PO 4: Conduct investigations of complex problems: Use research-based knowledge and research methods including design of experiments, analysis and interpretation of data, and synthesis of the information to provide valid conclusions.

PO 5: Modern tool usage: Create, select, and apply appropriate techniques, resources, and modern engineering and IT tools including prediction and modeling to complex engineering activities with an understanding of the limitations.

PO 6: The engineer and society: Apply reasoning informed by the contextual knowledge to assess societal, health, safety, legal and cultural issues and the consequent responsibilities relevant to the professional engineering practice.

PO 7: Environment and sustainability: Understand the impact of the professional engineering solutions in societal and environmental contexts, and demonstrate the knowledge of, and need for sustainable development.

PO 8: Ethics: Apply ethical principles and commit to professional ethics and responsibilities and norms of the engineering practice.

PO 9: Individual and team work: Function effectively as an individual, and as a member or leader in diverse teams, and in multidisciplinary settings.

PO 10: Communication: Communicate effectively on complex engineering activities with the engineering community and with society at large, such as, being able to comprehend and write effective reports and design documentation, make effective presentations, and give and receive clear instructions.

PO 11: Project management and finance: Demonstrate knowledge and understanding of the engineering and management principles and apply these to one's own work, as a member and leader in a team, to manage projects and in multidisciplinary environments.

PO 12: Life-long learning: Recognize the need for, and have the preparation and ability to engage in independent and life-long learning in the broadest context of technological change.

PROGRAM SPECIFIC OUTCOMES (PSOs)

PSO 1: To identify and define the computing requirements appropriate for its solution under given constraints.

PSO 2: To follow the best practices, namely, SEI-CMM levels and 6-sigma which varies from time to time for software development projects using open-ended programming environments to produce software deliverables as per customer needs.

6. Course Objectives and Course Outcomes

Course Objectives:

Develop ability to

1. Know the basic elements of Big Data and Data science to handle huge amount of data.
2. Gain knowledge of basic mathematics behind the Big data.
3. Understand the different Big data processing technologies.
4. Apply the Analytical concepts of Big data using R and Python.
5. Visualize the Big Data using different tools.

Course Outcomes (COs):

At the end of the course, student would be able to:

- CO1.** Observe Big Data elements and Architectures.
- CO2.** Apply different mathematical models for Big Data.
- CO3.** Demonstrate their Big Data skills by developing different applications.
- CO4.** Apply each learning model for different datasets.
- CO5.** Analyze needs, challenges and techniques for big data visualization.

Mapping of Lab Course with Programme Educational Objectives

| Course | Course Code | PEOs | POs & PSOs |
|---------------------------|-----------------|-------------------------|--|
| DATA ANALYTICS LAB | 18CS41L1 | PEO1, PEO2, PEO3 | PO1, PO2, PO3, PO4, PO5, PO11, PO12, PSO1, PSO2 |

Mapping of Lab Course outcomes with Programme outcomes:

| Course Outcomes - DATA ANALYTICS (18CS41L1) | Program Outcomes and Program Specific Outcomes | | | | | | | | | | | | | |
|--|---|---|---|---|---|---|---|---|---|----|----|----|------|------|
| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | PSO1 | PSO2 |
| CO1. Observe Big Data elements and Architectures. | 2 | 1 | 1 | 1 | 1 | - | - | - | - | - | 2 | 2 | 1 | 2 |
| CO2. Apply different mathematical models for Big Data. | 1 | 1 | 2 | 3 | 2 | - | - | - | - | - | 2 | 2 | 1 | 2 |
| CO3. Demonstrate their Big Data skills by developing different applications. | 2 | 1 | 1 | 2 | 1 | - | - | - | - | - | 1 | 2 | 1 | 2 |
| CO4. Apply each learning model for different datasets. | 1 | 1 | 1 | 2 | 1 | - | - | - | - | - | 1 | 2 | 1 | 2 |
| CO5. Analyze needs, challenges and techniques for big data visualization | 2 | 1 | 1 | 2 | 1 | - | - | - | - | - | 1 | 2 | 1 | 2 |

Prerequisites:

- 18CS2102 - Object Oriented Programming using Java
- 18CS2203 -Database Management Systems

INSTRUCTIONS TO THE STUDENTS:

1. Students are required to attend all labs.
2. Students should be dressed in formals when attending the laboratory sessions.
3. Students will work individually in computer laboratories.
4. While coming to the lab bring the observation book and Work book etc.
5. Before coming to the lab, prepare the pre-lab questions. Read through the lab experiment to familiarize you.
6. Utilize 3 hours' time properly to perform the experiment and noting down the outputs.
7. If the experiment is not completed in the prescribed time, the pending work has to be done in the leisure hour or extended hours.
8. You will be expected to submit the completed work book according to the deadlines set up by your instructor.

INSTRUCTIONS TO LABORATORY TEACHERS:

- Observation book and lab records submitted for the lab work are to be checked and signed before the next lab session.
- Students should be instructed to switch ON the power supply after the connections are checked by the lab assistant / teacher.
- The promptness of submission should be strictly insisted by awarding the marks accordingly.
- Ask viva questions at the end of the experiment.
- Do not allow students who come late to the lab class.
- Encourage the students to do the experiments innovatively.
- Fill continuous Evaluation sheet, on regular basis.
- Ensure that the students are dressed in formals

Scheme of Lab Exam Evaluation:

Evaluation of Internal Marks:

a) 15 Marks are awarded for day-to-day work

- 1) Record and Observation book ----- 5Marks
- 2) Attendance and behavior of student ----- 5 Marks
- 3) Viva and performance -----5 Marks

b) 15 Marks are awarded for conducting laboratory test as follows:

- 1) Write up and program-----5 Marks
- 2) Execution of Program -----5 Marks
- 3) Viva and performance -----5 Marks

Evaluation of External Marks:

70 Marks are awarded for conducting laboratory test as follows:

- 1) Algorithm ----- 25 Marks.
- 2) Write up and program----- 15 Marks
- 3) Execution of Program ----- 15 Marks
- 4) Viva ----- 15 Marks

PERFORMANCE INDICATOR

[illegible]

WEEK 1

INSTALLATION, CONFIGURATION AND RUNNING HADOOP AND HDFS.

PROCEDURE:

1. Installing Java

prince@prince-VirtualBox:~\$ cd ~

Update the source list

prince@prince-VirtualBox:~\$ sudo apt-get update

The OpenJDK project is the default version of Java that is provided from a supported Ubuntu repository.

prince@prince-VirtualBox:~\$ sudo apt-get install default-jdk

prince@prince-VirtualBox:~\$ java -version

java version "1.7.0_65"

OpenJDK Runtime Environment (IcedTea 2.5.3) (7u71-2.5.3-0ubuntu0.14.04.1)

OpenJDK 64-Bit Server VM (build 24.65-b04, mixed mode)

2. Adding a dedicated Hadoop user

prince@prince-VirtualBox:~\$ sudo add group hadoop

Adding group `hadoop' (GID 1002) ...
Done.

prince@prince-VirtualBox:~\$ sudo add user --ingroup hadoop hduser

Adding user `hduser' ...

Adding new user `hduser' (1001) with group `hadoop' ...

Creating home directory `/home/hduser' ...

Copying files from `/etc/skel' ...

Enter new UNIX password:

Retype new UNIX password:

passwd: password updated successfully

Changing the user information for hduser

Enter the new value, or press ENTER for the default

Full Name []:

Room Number []:

Work Phone []:

Home Phone []:

Other []:

Is the information correct? [Y/n] Y

3. Installing SSH

ssh has two main components:

1. **ssh** : The command we use to connect to remote machines - the client.
2. **sshd** : The daemon that is running on the server and allows clients to connect to the server.

The **ssh** is pre-enabled on Linux, but in order to start **sshd** daemon, we need to install **ssh** first. Use this command to do that :

```
prince@prince-VirtualBox:~$ sudo apt-get install ssh
```

This will install ssh on our machine. If we get something similar to the following, we can think it is setup properly:

```
prince@prince-VirtualBox:~$ which ssh
```

```
/usr/bin/ssh
```

```
prince@prince-VirtualBox:~$ which sshd
```

```
/usr/sbin/sshd
```

4. Create and Setup SSH Certificates

Hadoop requires SSH access to manage its nodes, i.e. remote machines plus our local

machine. For our single-node setup of Hadoop, we therefore need to configure SSH access to localhost.

So, we need to have SSH up and running on our machine and configured it to allow SSH public key authentication.

Hadoop uses SSH (to access its nodes) which would normally require the user to enter a password. However, this requirement can be eliminated by creating and setting up SSH certificates using the following commands. If asked for a filename just leave it blank and press the enter key to continue.

```
prince@prince-VirtualBox:~$ suhduser
```

Password:

```
prince@prince-VirtualBox:~$ ssh-keygen -t rsa -P ""
```

Generating public/private rsa key pair.

Enter file in which to save the key (/home/hduser/.ssh/id_rsa):

Created directory '/home/hduser/.ssh'.

Your identification has been saved in /home/hduser/.ssh/id_rsa.

Your public key has been saved in /home/hduser/.ssh/id_rsa.pub.

The key fingerprint is:

50:6b:f3:fc:0f:32:bf:30:79:c2:41:71:26:cc:7d:e3

hduser@prince-VirtualBox

The key's randomart image is:

```
+--[ RSA 2048]-----+
|  .oo.o  |
|  .o=. o  |
|  .+. o.  |
|  o = E  |
|  S +    |
|  . +    |
|  O +    |
|  O o    |
|  o..    |
+-----+
```

**hduser@prince-VirtualBox:/home/k\$ cat \$HOME/.ssh/id_rsa.pub >>
\$HOME/.ssh/authorized_keys**

The second command adds the newly created key to the list of authorized keys so that Hadoop can use ssh without prompting for a password.

We can check if ssh works:

hduser@prince-VirtualBox:/home/k\$ sshlocalhost

```
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ECDSA key fingerprint is e1:8b:a0:a5:75:ef:f4:b4:5e:a9:ed:be:64:be:5c:2f.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'localhost' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 14.04.1 LTS (GNU/Linux 3.13.0-40-generic x86_64)
...
```

5. Install Hadoop

**hduser@prince-VirtualBox:~\$ wget
http://mirrors.sonic.net/apache/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz
hduser@prince-VirtualBox:~\$ tar xvzf hadoop-2.6.0.tar.gz**

We want to move the Hadoop installation to the **/usr/local/hadoop** directory using the following command:

hduser@prince-VirtualBox:~/hadoop-2.6.0\$ sudo mv * /usr/local/hadoop

[sudo] password for hduser:
hduser is not in the sudoers file. This incident will be reported.

Oops!... We got:

"hduser is not in the sudoers file. This incident will be reported."

This error can be resolved by logging in as a root user, and then add **hduser** to **sudo**:

hduser@prince-VirtualBox:~/hadoop-2.6.0\$ su prince

Password:

prince@prince-VirtualBox:/home/hduser\$ sudo add user hduser sudo

[sudo] password for prince:

Adding user `hduser' to group `sudo' ...

Adding user hduser to group sudo

Done.

Now, the **hduser** has root privilege, we can move the Hadoop installation to the

/usr/local/hadoop directory without any problem:

prince@prince-VirtualBox:/home/hduser\$ sudosuhduser

hduser@prince-VirtualBox:~/hadoop-2.6.0\$ sudo mv * /usr/local/hadoop

hduser@prince-VirtualBox:~/hadoop-2.6.0\$ sudo chown -R hduser:hadoop /usr/local/hadoop

6. Setup Configuration Files

The following files will have to be modified to complete the Hadoop setup:

i. ~/.bashrc

ii. /usr/local/hadoop/etc/hadoop/hadoop-env.sh

iii. /usr/local/hadoop/etc/hadoop/core-site.xml

iv. /usr/local/hadoop/etc/hadoop/mapred-site.xml.template

v. /usr/local/hadoop/etc/hadoop/hdfs-site.xml

i. ~/.bashrc:

Before editing the **.bashrc** file in our home directory, we need to find the path where Java has been installed to set the **JAVA_HOME** environment variable using the following command:

hduser@prince-VirtualBox:~\$ update-alternatives --config java

There is only one alternative in link group java (providing /usr/bin/java): /usr/lib/jvm/java-7-openjdk-amd64/jre/bin/java

Nothing to configure.

Now we can append the following to the end of **~/.bashrc**:

hduser@prince-VirtualBox:~\$ nano ~/.bashrc

```
#HADOOP VARIABLES START
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
export HADOOP_INSTALL=/usr/local/hadoop
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"
#HADOOP VARIABLES END
```

hduser@prince-VirtualBox:~\$ source ~/.bashrc

note that the JAVA_HOME should be set as the path just before the '.../bin/':

hduser@ubuntu-VirtualBox:~\$ javac -version

javac 1.7.0_75

hduser@ubuntu-VirtualBox:~\$ which javac

/usr/bin/javac

hduser@ubuntu-VirtualBox:~\$ readlink -f /usr/bin/javac

/usr/lib/jvm/java-7-openjdk-amd64/bin/javac

ii. /usr/local/hadoop/etc/hadoop/hadoop-env.sh

We need to set **JAVA_HOME** by modifying **hadoop-env.sh** file.

hduser@prince-VirtualBox:~\$ nano /usr/local/hadoop/etc/hadoop/hadoop-env.sh

```
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
```

Adding the above statement in the **hadoop-env.sh** file ensures that the value of **JAVA_HOME** variable will be available to Hadoop whenever it is started up.

iii. /usr/local/hadoop/etc/hadoop/core-site.xml:

The **/usr/local/hadoop/etc/hadoop/core-site.xml** file contains configuration properties that Hadoop uses when starting up.

This file can be used to override the default settings that Hadoop starts with.

hduser@prince-VirtualBox:~\$ sudo mkdir -p /app/hadoop/tmp

hduser@prince-VirtualBox:~\$ sudo chown hduser:hadoop /app/hadoop/tmp

Open the file and enter the following in between the **<configuration>** tag:

hduser@prince-VirtualBox:~\$ nano /usr/local/hadoop/etc/hadoop/core-site.xml

```

<configuration>
<property>
<name>hadoop.tmp.dir</name>
<value>/app/hadoop/tmp</value>
<description>A base for other temporary directories.</description>
</property>

<property>
<name>fs.default.name</name>
<value>hdfs://localhost:54310</value>
<description>The name of the default file system. A URI whose
scheme and authority determine the FileSystem implementation. The
uri's scheme determines the config property (fs.SCHEME.impl) naming
theFileSystem implementation class. The uri's authority is used to
determine the host, port, etc. for a filesystem.</description>
</property>
</configuration>

```

iv. /usr/local/hadoop/etc/hadoop/mapred-site.xml

By default, the **/usr/local/hadoop/etc/hadoop/** folder contains **/usr/local/hadoop/etc/hadoop/mapred-site.xml.template** file which has to be renamed/copied with the name **mapred-site.xml**:

hduser@prince-VirtualBox:~\$ cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.template /usr/local/hadoop/etc/hadoop/mapred-site.xml

The **mapred-site.xml** file is used to specify which framework is being used for MapReduce. We need to enter the following content in between the **<configuration></configuration>** tag:

```

<configuration>
<property>
<name>mapred.job.tracker</name>
<value>localhost:54311</value>
<description>The host and port that the MapReduce job tracker runs
at. If "local", then jobs are run in-process as a single map
and reduce task.
</description>
</property>
</configuration>

```

v. /usr/local/hadoop/etc/hadoop/hdfs-site.xml

The **/usr/local/hadoop/etc/hadoop/hdfs-site.xml** file needs to be configured for each host in the cluster that is being used.

It is used to specify the directories which will be used as the **namenode** and the **datanode** on that host.

Before editing this file, we need to create two directories which will contain the namenode and the datanode for this Hadoop installation.

This can be done using the following commands:

```
hduser@prince-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop_store/hdfs/namenode
hduser@prince-VirtualBox:~$ sudo mkdir -p /usr/local/hadoop_store/hdfs/datanode
hduser@prince-VirtualBox:~$ sudo chown -R hduser:hadoop /usr/local/hadoop_store
```

Open the file and enter the following content in between the <configuration></configuration> tag:

```
hduser@prince-VirtualBox:~$ nano /usr/local/hadoop/etc/hadoop/hdfs-site.xml
```

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
<description>Default block replication.
  The actual number of replications can be specified when the file is created.
  The default is used if replication is not specified in create time.
</description>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:/usr/local/hadoop_store/hdfs/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>file:/usr/local/hadoop_store/hdfs/datanode</value>
</property>
</configuration>
```

7. Format the NewHadoop File system

Now, the Hadoop file system needs to be formatted so that we can start to use it. The format command should be issued with write permission since it creates **current** directory under **/usr/local/hadoop_store/hdfs/namenode** folder:

```
hduser@prince-VirtualBox:~$ hadoop namenode -format
```

DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

```
/*****
SHUTDOWN_MSG: Shutting down NameNode at laptop/192.168.1.1
*****/
```

Note that **hadoopnamenode -format** command should be executed once before we start using Hadoop.

If this command is executed again after Hadoop has been used, it'll destroy all the data on the Hadoop file system.

8. Starting Hadoop

Now it's time to start the newly installed single node cluster.
We can use **start-all.sh** or (**start-dfs.sh** and **start-yarn.sh**)

```
prince@prince-VirtualBox:~$ cd /usr/local/hadoop/sbin
```

```
prince@prince-VirtualBox:/usr/local/hadoop/sbin$ ls
```

```
distribute-exclude.sh  start-all.cmd      stop-balancer.sh
hadoop-daemon.sh       start-all.sh       stop-dfs.cmd
hadoop-daemons.sh     start-balancer.sh   stop-dfs.sh
hdfs-config.cmd        start-dfs.cmd       stop-secure-dns.sh
hdfs-config.sh         start-dfs.sh        stop-yarn.cmd
httpfs.sh              start-secure-dns.sh stop-yarn.sh
kms.sh                 start-yarn.cmd      yarn-daemon.sh
mr-jobhistory-daemon.sh start-yarn.sh       yarn-daemons.sh
refresh-namenodes.sh   stop-all.cmd
slaves.sh              stop-all.sh
```

```
prince@prince-VirtualBox:/usr/local/hadoop/sbin$ sudo su hduser
```

```
hduser@prince-VirtualBox:/usr/local/hadoop/sbin$ start-all.sh
```

```
hduser@prince-VirtualBox:~$ start-all.sh
```

This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh

```
15/04/18 16:43:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
```

```
Starting namenodes on [localhost]
```

```
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hduser-namenode-
laptop.out
```

```
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hduser-datanode-
laptop.out
```

```
Starting secondary namenodes [0.0.0.0]
```

```
0.0.0.0: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hduser-
secondarynamenode-laptop.out
```

```
15/04/18 16:43:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
```

```
starting yarn daemons
```

```
startingresourcemanager, logging to /usr/local/hadoop/logs/yarn-hduser-resourcemanager-
laptop.out
```

```
localhost: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hduser-nodemanager-
laptop.out
```

We can check if it's really up and running:

hduser@prince-VirtualBox:/usr/local/hadoop/sbin\$ jps

```
9026 NodeManager
7348 NameNode
9766 Jps
8887 ResourceManager
7507 DataNode
```

The output means that we now have a functional instance of Hadoop running on our VPS (Virtual private server).

9. Stopping Hadoop

\$ pwd

```
/usr/local/hadoop/sbin
```

hduser@prince-VirtualBox:/usr/local/hadoop/sbin\$ stop-all.sh

```
This script is Deprecated. Instead use stop-dfs.sh and stop-yarn.sh
15/04/18 15:46:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
Stopping namenodes on [localhost]
localhost: stopping namenode
localhost: stopping datanode
Stopping secondary namenodes [0.0.0.0]
0.0.0.0: no secondarynamenode to stop
15/04/18 15:46:59 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
stopping yarn daemons
stoppingresourcemanager
localhost: stopping nodemanager
noproxyserver to stop
```

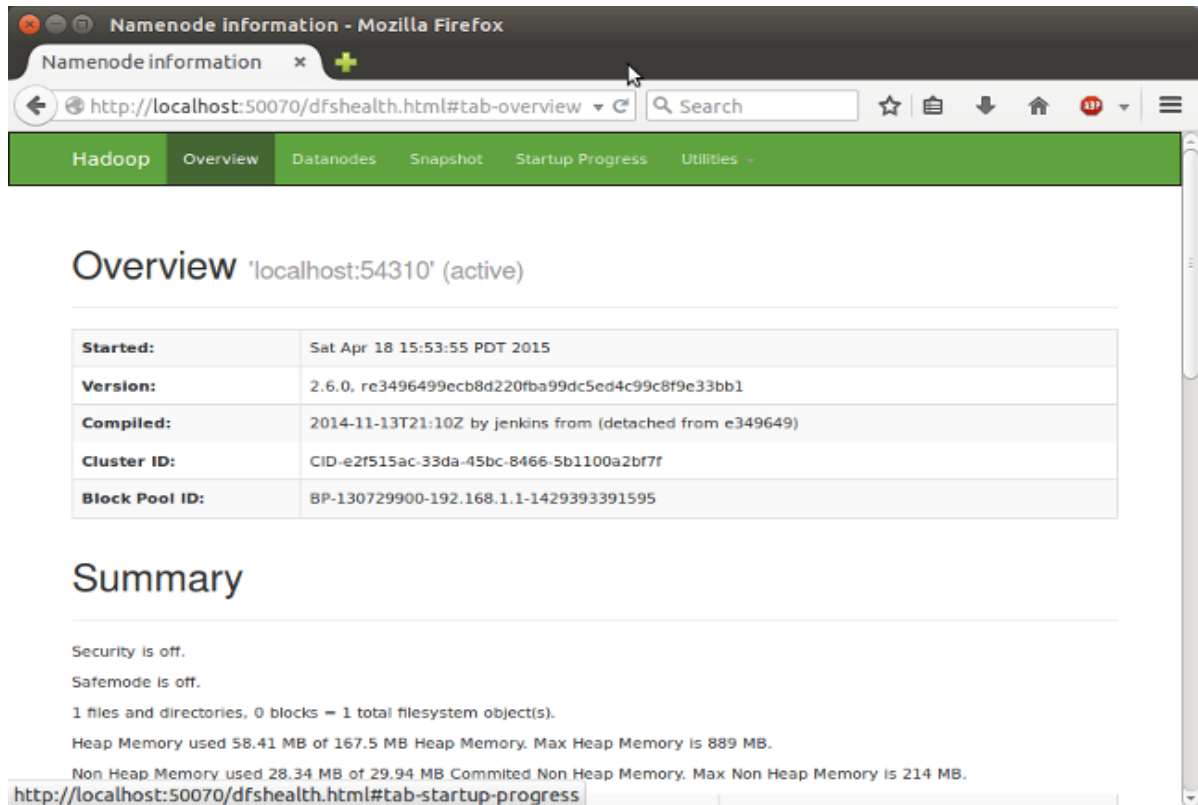
10.Hadoop Web Interfaces

Let's start the Hadoop again and see its Web UI:

```
hduser@prince-VirtualBox:/usr/local/hadoop/sbin$ start-all.sh
```

<http://localhost:50070/> - web UI of the NameNode daemon

OUTPUT:



Overview 'localhost:54310' (active)

| | |
|-----------------------|---|
| Started: | Sat Apr 18 15:53:55 PDT 2015 |
| Version: | 2.6.0, re3496499ecb8d220fba99dc5ed4c99c8f9e33bb1 |
| Compiled: | 2014-11-13T21:10Z by jenkins from (detached from e349649) |
| Cluster ID: | CID-e2f515ac-33da-45bc-8466-5b1100a2bf7f |
| Block Pool ID: | BP-130729900-192.168.1.1-1429393391595 |

Summary

Security is off.
Safemode is off.
1 files and directories, 0 blocks = 1 total filesystem object(s).
Heap Memory used 58.41 MB of 167.5 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 28.34 MB of 29.94 MB Committed Non Heap Memory. Max Non Heap Memory is 214 MB.

<http://localhost:50070/dfshealth.html#tab-startup-progress>

Namenode information - Mozilla Firefox

Namenode information

http://localhost:50070/dfshealth.html#tab-overview

Summary

Security is off.
Safemode is off.
1 files and directories, 0 blocks = 1 total filesystem object(s).
Heap Memory used 58.41 MB of 167.5 MB Heap Memory. Max Heap Memory is 889 MB.
Non Heap Memory used 28.34 MB of 29.94 MB Committed Non Heap Memory. Max Non Heap Memory is 214 MB.

| | |
|---|-------------------------------|
| Configured Capacity: | 454.29 GB |
| DFS Used: | 24 KB |
| Non DFS Used: | 125.8 GB |
| DFS Remaining: | 328.49 GB |
| DFS Used%: | 0% |
| DFS Remaining%: | 72.31% |
| Block Pool Used: | 24 KB |
| Block Pool Used%: | 0% |
| DataNodes usages% (Min/Median/Max/stdDev): | 0.00% / 0.00% / 0.00% / 0.00% |
| Live Nodes | 1 (Decommissioned: 0) |
| Dead Nodes | 0 (Decommissioned: 0) |
| Decommissioning Nodes | 0 |
| Number of Under-Replicated Blocks | 0 |
| Number of Blocks Pending Deletion | 0 |
| Block Deletion Start Time | 4/18/2015, 3:53:55 PM |

Namenode information - Mozilla Firefox

Namenode information

http://localhost:50070/dfshealth.html#tab-overview

| | |
|---|-------------------------------|
| DataNodes usages% (Min/Median/Max/stdDev): | 0.00% / 0.00% / 0.00% / 0.00% |
| Live Nodes | 1 (Decommissioned: 0) |
| Dead Nodes | 0 (Decommissioned: 0) |
| Decommissioning Nodes | 0 |
| Number of Under-Replicated Blocks | 0 |
| Number of Blocks Pending Deletion | 0 |
| Block Deletion Start Time | 4/18/2015, 3:53:55 PM |

NameNode Journal Status

Current transaction ID: 2

| Journal Manager | State |
|--|---|
| FileJournalManager(root=/usr/local/hadoop_store/hdfs/namenode) | EditLogFileOutputStream(/usr/local/hadoop_store/hdfs/namenode/current/edits_inprogress_0000000000000000002) |

NameNode Storage


| Storage Directory | Type | State |
|---------------------------------------|-----------------|--------|
| /usr/local/hadoop_store/hdfs/namenode | IMAGE_AND_EDITS | Active |

Hadoop, 2014.

Legacy UI

NameNode Journals

Secondary Name Node



SecondaryNameNode

| | |
|------------------|---|
| Version: | 2.6.0, e3496499ecb8d220fba99dc5ed4c99c8f9e33bb1 |
| Compiled: | 2014-11-13T21:10Z by jenkins from (detached from e349649) |

SecondaryNameNode Status

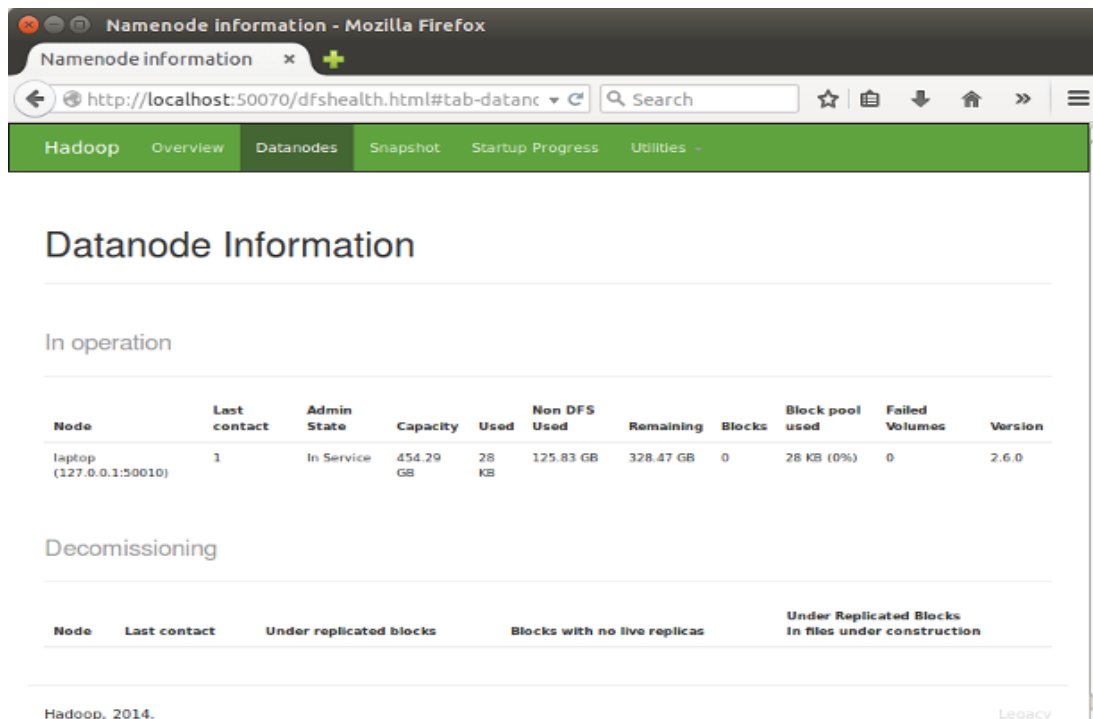
Name Node Address : localhost/127.0.0.1:54310
Start Time : Sat Apr 18 16:43:38 PDT 2015
Last Checkpoint : 79 seconds ago
Checkpoint Period : 3600 seconds
Checkpoint Transactions: 1000000
Checkpoint Dirs : [file:///app/hadoop/tmp/dfs/namesecondary]
Checkpoint Edits Dirs : [file:///app/hadoop/tmp/dfs/namesecondary]

[Logs](#)

[Hadoop](#), 2015.

(Note) I had to restart Hadoop to get this Secondary Namenode.


Data Node



| Node | Last contact | Admin State | Capacity | Used | Non DFS Used | Remaining | Blocks | Block pool used | Failed Volumes | Version |
|--------------------------|--------------|-------------|-----------|-------|--------------|-----------|--------|-----------------|----------------|---------|
| laptop (127.0.0.1:50010) | 1 | In Service | 454.29 GB | 28 KB | 125.83 GB | 328.47 GB | 0 | 28 KB (0%) | 0 | 2.6.0 |

| Node | Last contact | Under replicated blocks | Blocks with no live replicas | Under Replicated Blocks in files under construction |
|------|--------------|-------------------------|------------------------------|---|
|------|--------------|-------------------------|------------------------------|---|

Hadoop, 2014. Legacy



Directory: /logs/

| | | |
|--|--------------|-------------------------|
| SecurityAuth-hduser.audit | 0 bytes | Apr 18, 2015 3:40:58 PM |
| hadoop-hduser-datanode-laptop.log | 72879 bytes | Apr 18, 2015 4:44:13 PM |
| hadoop-hduser-datanode-laptop.out | 718 bytes | Apr 18, 2015 4:43:21 PM |
| hadoop-hduser-datanode-laptop.out.1 | 718 bytes | Apr 18, 2015 3:53:49 PM |
| hadoop-hduser-datanode-laptop.out.2 | 718 bytes | Apr 18, 2015 3:41:03 PM |
| hadoop-hduser-namenode-laptop.log | 121216 bytes | Apr 18, 2015 4:52:23 PM |
| hadoop-hduser-namenode-laptop.out | 718 bytes | Apr 18, 2015 4:43:16 PM |
| hadoop-hduser-namenode-laptop.out.1 | 718 bytes | Apr 18, 2015 3:53:44 PM |
| hadoop-hduser-namenode-laptop.out.2 | 718 bytes | Apr 18, 2015 3:40:58 PM |
| hadoop-hduser-secondarynamenode-laptop.log | 51913 bytes | Apr 18, 2015 4:52:38 PM |
| hadoop-hduser-secondarynamenode-laptop.out | 718 bytes | Apr 18, 2015 4:43:37 PM |
| hadoop-hduser-secondarynamenode-laptop.out.1 | 718 bytes | Apr 18, 2015 3:54:06 PM |
| hadoop-hduser-secondarynamenode-laptop.out.2 | 718 bytes | Apr 18, 2015 3:42:52 PM |
| userlogs/ | 4096 bytes | Apr 18, 2015 4:52:22 PM |
| yarn-hduser-nodemanager-laptop.log | 81625 bytes | Apr 18, 2015 4:44:32 PM |
| yarn-hduser-nodemanager-laptop.out | 702 bytes | Apr 18, 2015 4:44:02 PM |
| yarn-hduser-nodemanager-laptop.out.1 | 702 bytes | Apr 18, 2015 3:54:32 PM |
| yarn-hduser-nodemanager-laptop.out.2 | 702 bytes | Apr 18, 2015 3:43:10 PM |
| yarn-hduser-resourcemanager-laptop.log | 107718 bytes | Apr 18, 2015 4:44:32 PM |
| yarn-hduser-resourcemanager-laptop.out | 702 bytes | Apr 18, 2015 4:44:00 PM |
| yarn-hduser-resourcemanager-laptop.out.1 | 702 bytes | Apr 18, 2015 3:54:29 PM |
| yarn-hduser-resourcemanager-laptop.out.2 | 702 bytes | Apr 18, 2015 3:43:08 PM |

VIVA QUESTIONS

1. What is Hadoop?

Apache Hadoop is an open-source framework that is used to efficiently store and process large datasets ranging in size from gigabytes to petabytes of data

2. What are the three modes in which Hadoop can run?

- Standalone Mode.
- Pseudo-distributed Mode.
- Fully-Distributed Mode.

3. What are the different Hadoop Configuration files?

HADOOP-ENV.sh, CORE-SITE.XML, HDFS-SITE.XML, MAPRED-SITE.XML, Masters and Slave

4. Difference between regular file system and HDFS?

Normal file systems have small block size of data. (Around 512 bytes) while HDFS has larger block sizes at around 64 MB) Multiple disks seek for larger files in normal file systems while in HDFS, data is read sequentially after every individual seek

5. What are the two main purpose of Hadoop?

Storing and processing big data

WEEK2:

IMPLEMENTATION OF WORD COUNT / FREQUENCY PROGRAMS USING MAPREDUCE

PROCEDURE:

1. Install hadoop.
2. Start all services using the command.

```
hduser@prince-VirtualBox:/usr/local/hadoop/bin$ jps
3242 Jps
```

```
hduser@prince-VirtualBox:/usr/local/hadoop/bin$ start-all.sh
```

```
hduser@prince-VirtualBox:/usr/local/hadoop/bin$ jps
16098 NameNode
16214 DataNode
16761 NodeManager
16636 ResourceManager
16429 SecondaryNameNode
19231 Jps
```

PROGRAM CODING:

```
hduser@prince-VirtualBox:/usr/local/hadoop/bin$ nano wordcount7.java
```

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class wordcount7 {

    public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable>
    {

        private final static IntWritable one = new IntWritable (1);
        private Text word = new Text ();

        public void map (Object key, Text value, Context context) throws IOException,
        InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
```

```

while (itr.hasMoreTokens ()) {
word.set (itr.nextToken ());
context.write (word, one);
}
}
}
public static class IntSumReducer
extends Reducer<Text,IntWritable,Text,IntWritable> {
privateIntWritable result = new IntWritable();

public void reduce(Text key, Iterable<IntWritable> values,
Context context
) throws IOException, InterruptedException {
int sum = 0;
for (IntWritableval : values) {
sum += val.get();
}
result.set(sum);
context.write(key, result);
}
}

public static void main(String[] args) throws Exception {
Configuration conf = new Configuration();
Job job = Job.getInstance (conf, "word count");
job.setJarByClass(wordcount7.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass (Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit (job.waitForCompletion (true) ? 0 : 1);
}
}

```

TO COMPILE:

hduser@prince-VirtualBox:/usr/local/hadoop/bin\$ **hadoopcom.sun.tools.javac.Main wordcount7.java**

TO CREATE A JAR FILE:

hduser@prince-VirtualBox:/usr/local/hadoop/bin\$ **jar cf wc2.jar wordcount7*.java**

TO CREATE A DIRECTORY IN HDFS:

hduser@prince-VirtualBox:/usr/local/hadoop/bin\$**hadoopdfs -mkdir /deepika**

TO LOAD INPUT FILE:

```
hduser@prince-VirtualBox:/usr/local/hadoop/bin$hdfs -put  
/home/prince/Downloads/wc.txt /deepika/wc1
```

TO EXECUTE:

```
hduser@prince-VirtualBox:/usr/local/hadoop/bin$ hadoop jar wc2.jar wordcount7  
/deepika/wc1.txt /deepika/out2
```

```
16/09/16 14:34:23 INFO mapreduce.Job: map 0% reduce 0%
```

```
.  
. .
```

```
16/09/16 14:34:26 INFO mapreduce.Job: map 100% reduce 0%
```

```
.  
. .
```

```
16/09/16 14:34:30 INFO mapreduce.Job: map 100% reduce 100%
```

```
.  
. .
```

```
16/09/16 14:34:31 INFO mapreduce.Job: Job job_local364071501_0001 completed  
successfully
```

```
16/09/16 14:34:31 INFO mapreduce.Job: Counters: 38
```

File System Counters

FILE: Number of bytes read=8552

FILE: Number of bytes written=507858

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=1424

HDFS: Number of bytes written=724

HDFS: Number of read operations=13

HDFS: Number of large read operations=0

HDFS: Number of write operations=4

Map-Reduce Framework

Map input records=10

Map output records=92

Map output bytes=1079

Map output materialized bytes=1018

Input split bytes=99

Combine input records=92

Combine output records=72

Reduce input groups=72

Reduce shuffle bytes=1018

Reduce input records=72

Reduce output records=72

Spilled Records=144

Shuffled Maps =1

Failed Shuffles=0

Merged Map outputs=1

GC time elapsed (ms)=111

CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=242360320
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=712
File Output Format Counters
Bytes Written=724

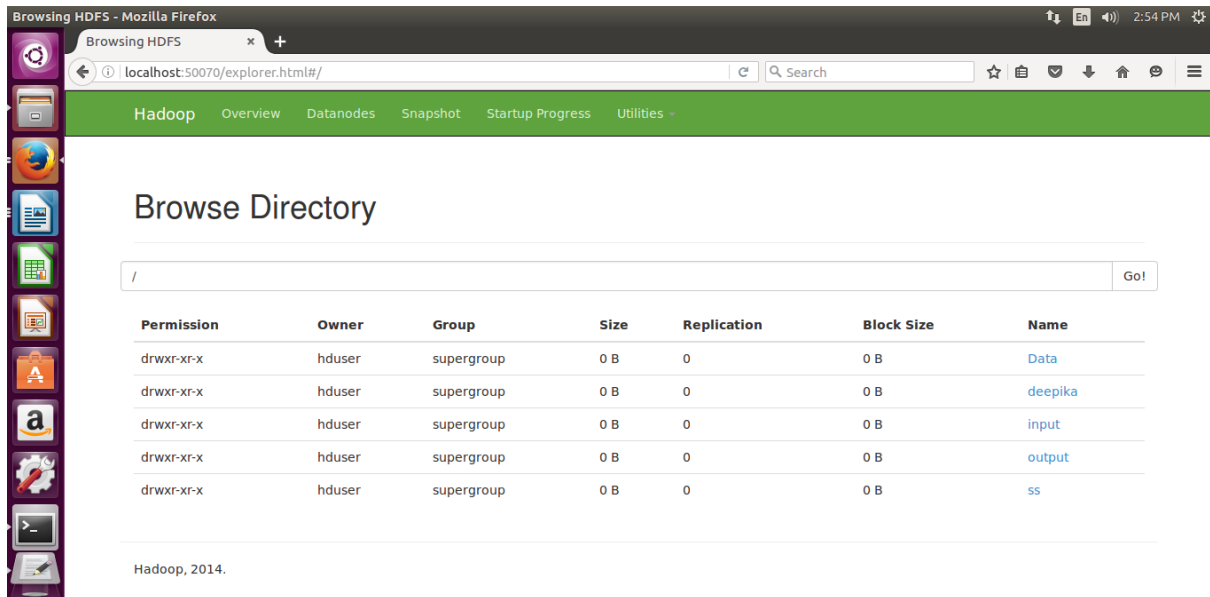
INPUT FILE:

wc1.txt

STEPS:

1. Open an editor and type WordCount program and save as WordCount.java
2. Set the path as export HADOOP_CLASSPATH=\${JAVA_HOME}/lib/tools.jar
3. To compile the program, bin/hadoopcom.sun.tools.javac.Main WordCount.java
4. Create a jar file, jar cf wc.jar WordCount*.class
5. Create input files input.txt,input1.txt and input2.txt and create a directory in hdfs, /mit/wordcount/input
6. Move these i/p files to hdfs system, bin/hadoopfs -put /opt/hadoop-2.7.0/input.txt /mit/wordcount/input/input.txt repeat this step for other two i/p files.
7. To execute, bin/hadoop jar wc.jar WordCount /mit/wordcount/input /mit/wordcount/output.
8. The mapreduce result will be available in the output directory.

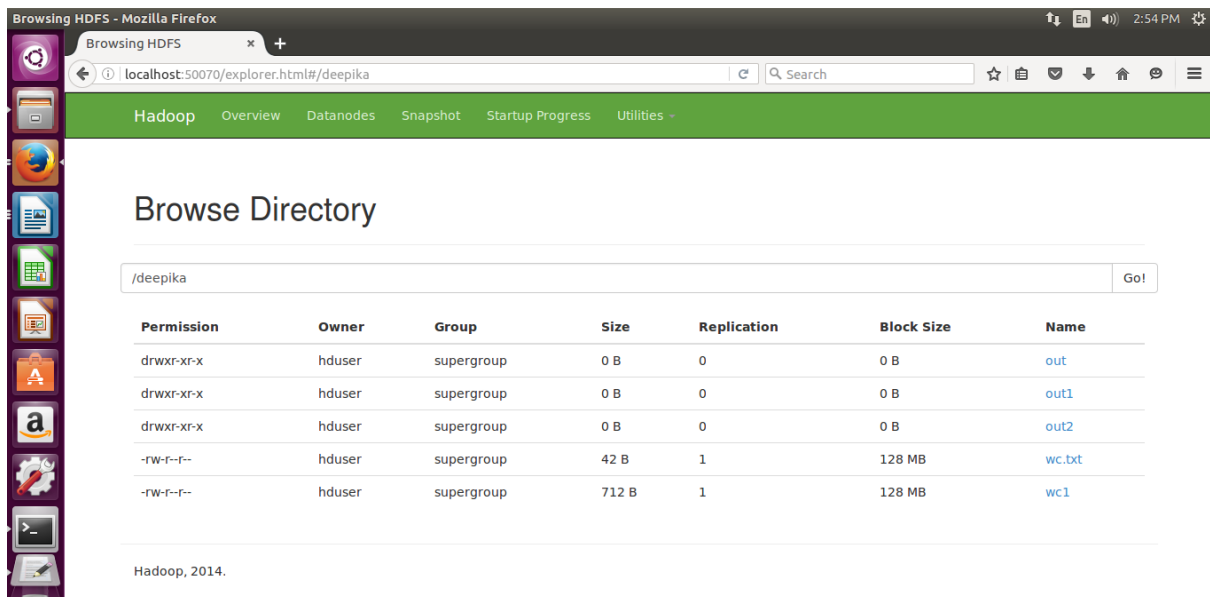
OUTPUT:



The screenshot shows the HDFS Explorer interface in a Mozilla Firefox browser. The address bar displays 'localhost:50070/explorer.html#/'. The navigation bar includes 'Hadoop', 'Overview', 'Datanodes', 'Snapshot', 'Startup Progress', and 'Utilities'. The main heading is 'Browse Directory'. Below it, a search bar contains the root path '/'. A table lists the contents of the root directory:

| Permission | Owner | Group | Size | Replication | Block Size | Name |
|------------|--------|------------|------|-------------|------------|-------------------------|
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | Data |
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | deepika |
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | input |
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | output |
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | ss |

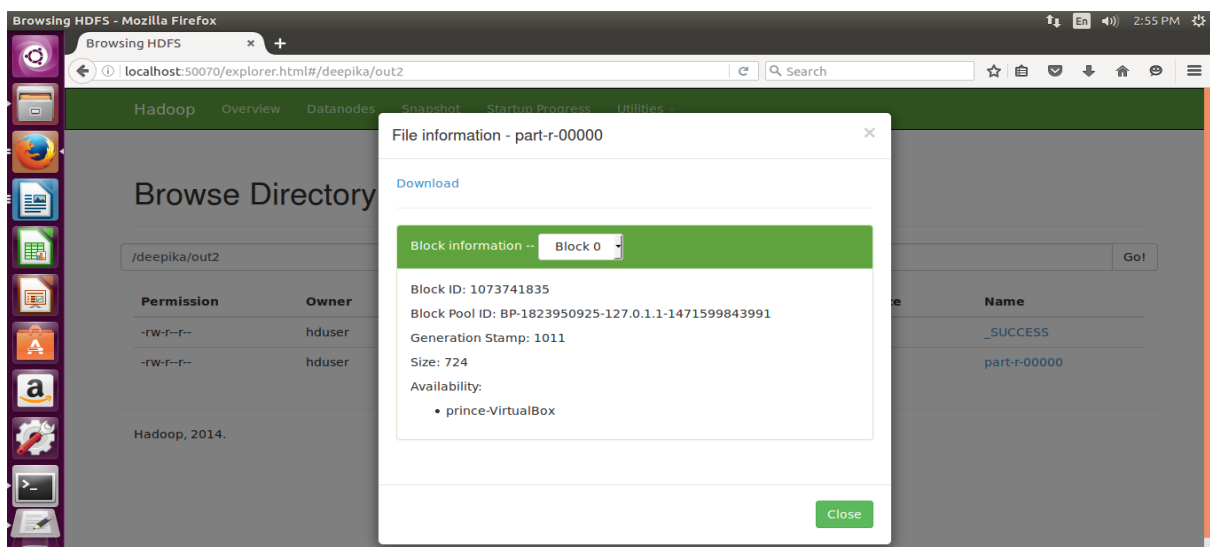
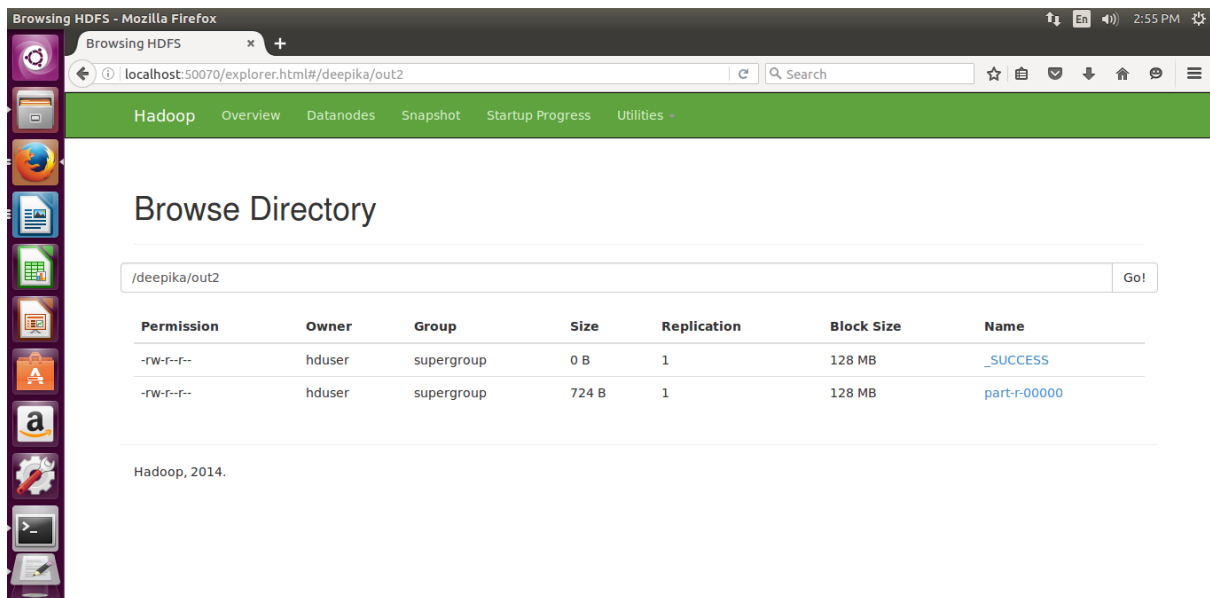
At the bottom, it says 'Hadoop, 2014.'



The screenshot shows the HDFS Explorer interface in a Mozilla Firefox browser. The address bar displays 'localhost:50070/explorer.html#/deepika'. The navigation bar is the same as the previous screenshot. The main heading is 'Browse Directory'. Below it, a search bar contains the path '/deepika'. A table lists the contents of the 'deepika' directory:

| Permission | Owner | Group | Size | Replication | Block Size | Name |
|------------|--------|------------|-------|-------------|------------|------------------------|
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | out |
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | out1 |
| drwxr-xr-x | hduser | supergroup | 0 B | 0 | 0 B | out2 |
| -rw-r--r-- | hduser | supergroup | 42 B | 1 | 128 MB | wc.txt |
| -rw-r--r-- | hduser | supergroup | 712 B | 1 | 128 MB | wc1 |

At the bottom, it says 'Hadoop, 2014.'



```
/mit/wordcount/input 2
/mit/wordcount/input/input.txt 1
/mit/wordcount/output. 1
/opt/hadoop-2.7.0/input.txt 1
1. 1
2. 1
3. 1
4. 1
5. 1
6. 1
7. 1
8. 1
```

Create 2

HADOOP_CLASSPATH=\${JAVA_HOME}/lib/tools.jar 1
Move 1
Open 1
STEPS: 1
Set 1
The 1
To 2
WordCount 2
WordCount*.class 1
WordCount.java 2
a 2
an 1
and 4
as 2
available 1
be 1
bin/hadoop 3
cf 1
com.sun.tools.javac.Main 1
compile 1
create 1
directory 1
directory. 1
editor 1
execute, 1
export 1
file, 1
files 2
files. 1
for 1
fs 1
hdfs 1
hdfs, 1
i/p 2
in 2
input 1
input.txt,input1.txt 1
input2.txt 1
jar 3
mapreduce 1
other 1
output 1
path 1
program 1
program, 1
repeat 1
result 1
save 1
step 1
system, 1

the 3
these 1
this 1
to 1
two 1
type 1
wc.jar 2
will 1
-put 1

Viva Questions

1. What is Map Reduce?

MapReduce implements various mathematical algorithms to divide a task into small parts and assign them to multiple systems. In technical terms, MapReduce algorithm helps in sending the Map & Reduce tasks to appropriate servers in a cluster.

2. Mention three advantages of Map Reduce?

Parallel processing, Scalability, Fast

3. What are the main components of Map Reduce?

The two main components of the MapReduce Job are the JobTracker and TaskTracker.

4. What is HDFS?

HDFS is a distributed file system that handles large data sets running on commodity hardware. It is used to scale a single Apache Hadoop cluster to hundreds (and even thousands) of nodes. HDFS is one of the major components of Apache Hadoop, the others being MapReduce and YARN.

5. What is meant by Job Tracker?

The JobTracker is the service within Hadoop that farms out MapReduce tasks to specific nodes in the cluster, ideally the nodes that have the data, or at least are in the same rack. Client applications submit jobs to the Job tracker. The JobTracker talks to the NameNode to determine the location of the data.

WEEK 3

Implementation of MR program that processes a weather dataset PROGRAM CODING:

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.output.MultipleOutputs;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

```
/**
```

```
 * @author devinline
```

```
*/
```

```
public class CalculateMaxAndMinTemperatureWithTime {  
public static String calOutputName = "California";  
public static String nyOutputName = "Newyork";  
public static String njOutputName = "Newjersy";  
public static String ausOutputName = "Austin";  
public static String bosOutputName = "Boston";  
public static String balOutputName = "Baltimore";
```

```
public static class WhetherForecastMapper extends
```

```
Mapper<Object, Text, Text, Text> {
```

```
public void map(Object keyOffset, Text dayReport, Context con)
```

```
    throws IOException, InterruptedException {
```

```
    StringTokenizer strTokens = new StringTokenizer(  
        dayReport.toString(), "\\t");
```

```
    int counter = 0;
```

```
    Float currnetTemp = null;
```

```
    Float minTemp = Float.MAX_VALUE;
```

```
    Float maxTemp = Float.MIN_VALUE;
```

```
    String date = null;
```

```
    String currentTime = null;
```

```
    String minTempANDTime = null;
```

```
    String maxTempANDTime = null;
```

```
while (strTokens.hasMoreElements()) {
```

```
    if (counter == 0) {
```

```
        date = strTokens.nextToken();
```

```
    } else {
```

```
        if (counter % 2 == 1) {
```

```
            currentTime = strTokens.nextToken();
```

```

    } else {
        currnetTemp = Float.parseFloat(strTokens.nextToken());
        if (minTemp > currnetTemp) {
            minTemp = currnetTemp;
            minTempANDTime = minTemp + "AND" + currentTime;
        }
        if (maxTemp < currnetTemp) {
            maxTemp = currnetTemp;
            maxTempANDTime = maxTemp + "AND" + currentTime;
        }
    }
    counter++;
}
// Write to context - MinTemp, MaxTemp and corresponding time
Text temp = new Text();
temp.set(maxTempANDTime);
Text dateText = new Text();
dateText.set(date);
try {
    con.write(dateText, temp);
} catch (Exception e) {
    e.printStackTrace();
}

temp.set(minTempANDTime);
dateText.set(date);
con.write(dateText, temp);

}
}

```

public static class WhetherForecastReducer extends

```

    Reducer<Text, Text, Text, Text> {
        MultipleOutputs<Text, Text> mos;

```

```

    public void setup(Context context) {
        mos = new MultipleOutputs<Text, Text>(context);
    }

```

```

    public void reduce(Text key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {
        int counter = 0;
        String reducerInputStr[] = null;
        String f1Time = "";
        String f2Time = "";
        String f1 = "", f2 = "";
        Text result = new Text();
        for (Text value : values) {

```

```

if (counter == 0) {
    reducerInputStr = value.toString().split("AND");
    f1 = reducerInputStr[0];
    f1Time = reducerInputStr[1];
}

else {
    reducerInputStr = value.toString().split("AND");
    f2 = reducerInputStr[0];
    f2Time = reducerInputStr[1];
}

counter = counter + 1;
}
if (Float.parseFloat(f1) > Float.parseFloat(f2)) {

    result = new Text("Time: " + f2Time + " MinTemp: " + f2 + "\t"
        + "Time: " + f1Time + " MaxTemp: " + f1);
} else {

    result = new Text("Time: " + f1Time + " MinTemp: " + f1 + "\t"
        + "Time: " + f2Time + " MaxTemp: " + f2);
}
String fileName = "";
if (key.toString().substring(0, 2).equals("CA")) {
    fileName = CalculateMaxAndMinTemperatureTime.calOutputName;
} else if (key.toString().substring(0, 2).equals("NY")) {
    fileName = CalculateMaxAndMinTemperatureTime.nyOutputName;
} else if (key.toString().substring(0, 2).equals("NJ")) {
    fileName = CalculateMaxAndMinTemperatureTime.njOutputName;
} else if (key.toString().substring(0, 3).equals("AUS")) {
    fileName = CalculateMaxAndMinTemperatureTime.ausOutputName;
} else if (key.toString().substring(0, 3).equals("BOS")) {
    fileName = CalculateMaxAndMinTemperatureTime.bosOutputName;
} else if (key.toString().substring(0, 3).equals("BAL")) {
    fileName = CalculateMaxAndMinTemperatureTime.balOutputName;
}
String strArr[] = key.toString().split("_");
key.set(strArr[1]); //Key is date value
mos.write(fileName, key, result);
}

```

@Override

```

public void cleanup(Context context) throws IOException,
    InterruptedException {
    mos.close();
}

```

```

public static void main(String[] args) throws IOException,

```

```

ClassNotFoundException, InterruptedException {
Configuration conf = new Configuration();
Job job = Job.getInstance(conf, "Wheather Statistics of USA");
job.setJarByClass(CalculateMaxAndMinTemperatureWithTime.class);

job.setMapperClass(WhetherForecastMapper.class);
job.setReducerClass(WhetherForecastReducer.class);

job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(Text.class);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(Text.class);

MultipleOutputs.addNamedOutput(job, calOutputName,
    TextOutputFormat.class, Text.class, Text.class);
MultipleOutputs.addNamedOutput(job, nyOutputName,
    TextOutputFormat.class, Text.class, Text.class);
MultipleOutputs.addNamedOutput(job, njOutputName,
    TextOutputFormat.class, Text.class, Text.class);
MultipleOutputs.addNamedOutput(job, bosOutputName,
    TextOutputFormat.class, Text.class, Text.class);
MultipleOutputs.addNamedOutput(job, ausOutputName,
    TextOutputFormat.class, Text.class, Text.class);
MultipleOutputs.addNamedOutput(job, balOutputName,
    TextOutputFormat.class, Text.class, Text.class);
// FileInputFormat.addInputPath(job, new Path(args[0]));
// FileOutputFormat.setOutputPath(job, new Path(args[1]));
Path pathInput = new Path(
    "hdfs://192.168.213.133:54310/weatherInputData/input_temp.txt");
Path pathOutputDir = new Path(
    "hdfs://192.168.213.133:54310/user/hduser1/testfs/output_mapred3");
FileInputFormat.addInputPath(job, pathInput);
FileOutputFormat.setOutputPath(job, pathOutputDir);
try {
    System.exit(job.waitForCompletion(true) ? 0 : 1);
} catch (Exception e) {
    // TODO Auto-generated catch block
    e.printStackTrace();
}
}
}

```

OUTPUT:

whether output directory is in place on HDFS. Execute following command to verify the same.

```

hduser1@ubuntu:/usr/local/hadoop2.6.1/bin$ ./hadoop fs -ls
/user/hduser1/testfs/output_mapred3

```

Found 8 items

```
-rw-r--r--  3 zytham supergroup    438 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/Austin-r-00000
-rw-r--r--  3 zytham supergroup    219 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/Baltimore-r-00000
-rw-r--r--  3 zytham supergroup    219 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/Boston-r-00000
-rw-r--r--  3 zytham supergroup    511 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/California-r-00000
-rw-r--r--  3 zytham supergroup    146 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/Newjersy-r-00000
-rw-r--r--  3 zytham supergroup    219 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/Newyork-r-00000
-rw-r--r--  3 zytham supergroup      0 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/_SUCCESS
-rw-r--r--  3 zytham supergroup      0 2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/part-r-00000
```

hduser1@ubuntu:/usr/local/hadoop2.6.1/bin\$./hadoop fs -cat

/user/hduser1/testfs/output_mapred3/Austin-r-00000

25-Jan-2018 Time: 12:34:542 MinTemp: -22.3 Time: 05:12:345 MaxTemp: 35.7

26-Jan-2018 Time: 22:00:093 MinTemp: -27.0 Time: 05:12:345 MaxTemp: 55.7

27-Jan-2018 Time: 02:34:542 MinTemp: -22.3 Time: 05:12:345 MaxTemp: 55.7

29-Jan-2018 Time: 14:00:093 MinTemp: -17.0 Time: 02:34:542 MaxTemp: 62.9

30-Jan-2018 Time: 22:00:093 MinTemp: -27.0 Time: 05:12:345 MaxTemp: 49.2

31-Jan-2018 Time: 14:00:093 MinTemp: -17.0 Time: 03:12:187 MaxTemp: 56.0

VIVA QUESTIONS

1. What is Text Input format?

TextInputFormat is one of the file formats of Hadoop. As the name suggest, it is used to read lines of text files. Basically, it helps in generating key-value pairs from the text.

2. What is difference between HDFS block and InputSplit?

Block – HDFS Block is the physical representation of data in Hadoop. InputSplit – MapReduce InputSplit is the logical representation of data present in the block in Hadoop. It is basically used during data processing in MapReduce program or other processing techniques.

3. What is partitioning?

Hadoop Partitioning specifies that all the values for each key are grouped together. It also makes sure that all the values of a single key go to the same reducer. This allows even distribution of the map output over the reducer.

4. How to set which framework would use to run mapreduce program?

MapReduce is a programming model or pattern within the Hadoop framework that is used to access big data stored in the Hadoop File System (HDFS). It is a core component, integral to the functioning of the Hadoop framework.

5. What is shuffling in map reduce?

The process of transferring data from the mappers to reducers is shuffling. It is also the process by which the system performs the sort. Then it transfers the map output to the reducer as input.

WEEK4

IMPLEMENTATION OF LINEAR AND LOGISTIC REGRESSION

PROGRAM CODING :(LINEAR REGRESSION)

The aim of linear regression is to model a continuous variable Y as a mathematical function of one or more X variable(s), so that we can use this regression model to predict the Y when only the X is known. This mathematical equation can be generalized as follows:

$$Y = \beta_1 + \beta_2 X + \epsilon$$

For this analysis, we will use the cars dataset that comes with R by default. cars is a standard built-in dataset, that makes it convenient to demonstrate linear regression in a simple and easy to understand fashion. You can access this dataset simply by typing in cars in your R console.

```
# display the first 6 observations
```

```
head(cars)
```

Scatter Plot:

```
# scatterplot
```

```
scatter.smooth(x=cars$speed, y=cars$dist, main="Dist ~ Speed")
```

Build Linear Model:

```
# build linear regression model on full data
```

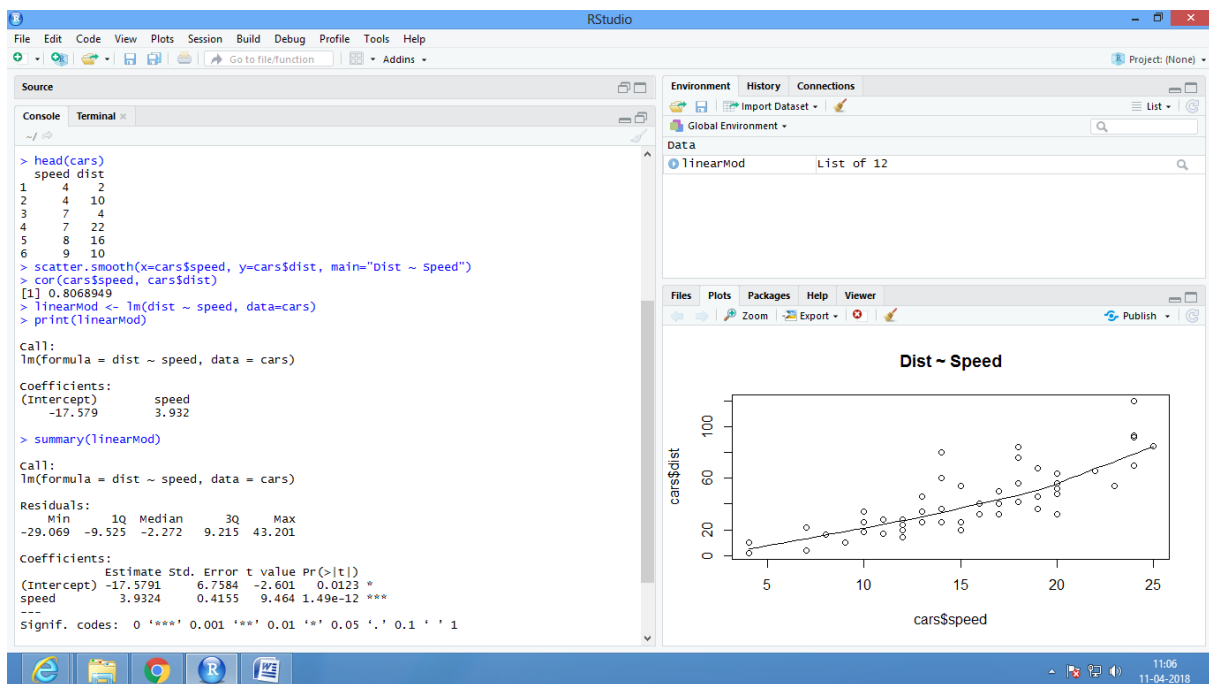
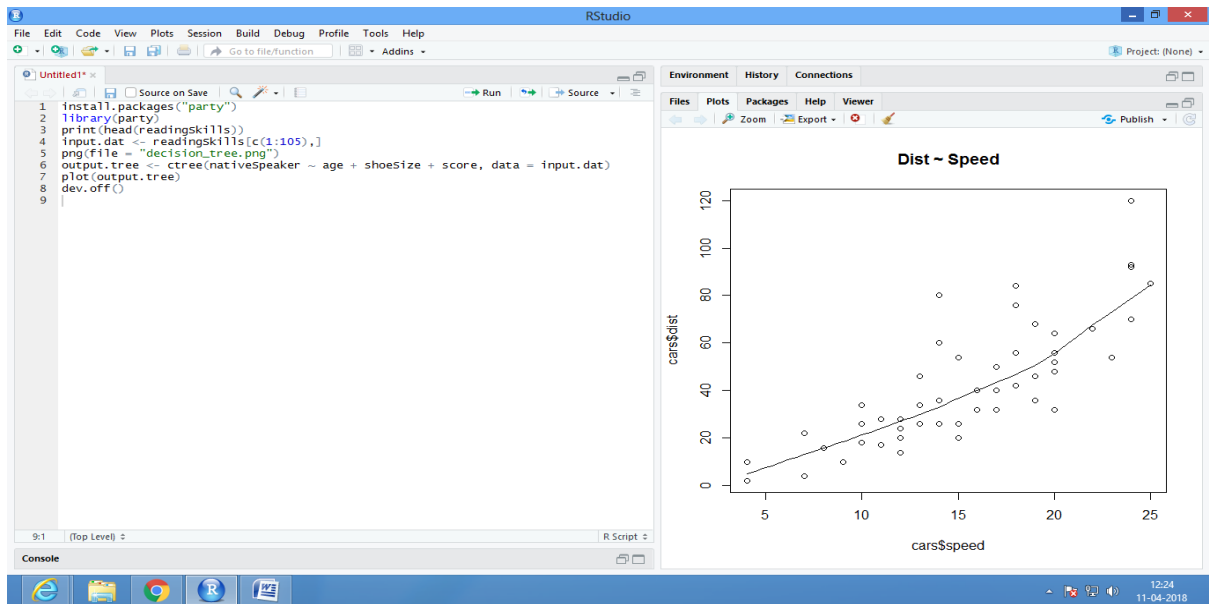
```
linearMod <- lm(dist ~ speed, data=cars)
```

```
print(linearMod)
```

```
Linear Regression Diagnostics
```

```
summary(linearMod)
```

OUTPUT:



PROGRAM CODING:(LOGISTIC REGRESSION):

The general mathematical equation for logistic regression is
$$y = 1/(1+e^{-(a+b_1x_1+b_2x_2+b_3x_3+\dots)})$$

The basic syntax for glm() function in logistic regression is
glm(formula,data,family)

The in-built data set "mtcars" describes different models of a car with their various engine specifications. In "mtcars" data set, the transmission mode (automatic or manual) is described by the column am which is a binary value (0 or 1). We can create a logistic regression model between the columns "am" and 3 other columns - hp, wt and cyl.

Select some columns form mtcars.

```
input <- mtcars[,c("am","cyl","hp","wt")]
```

```
print(head(input))
```

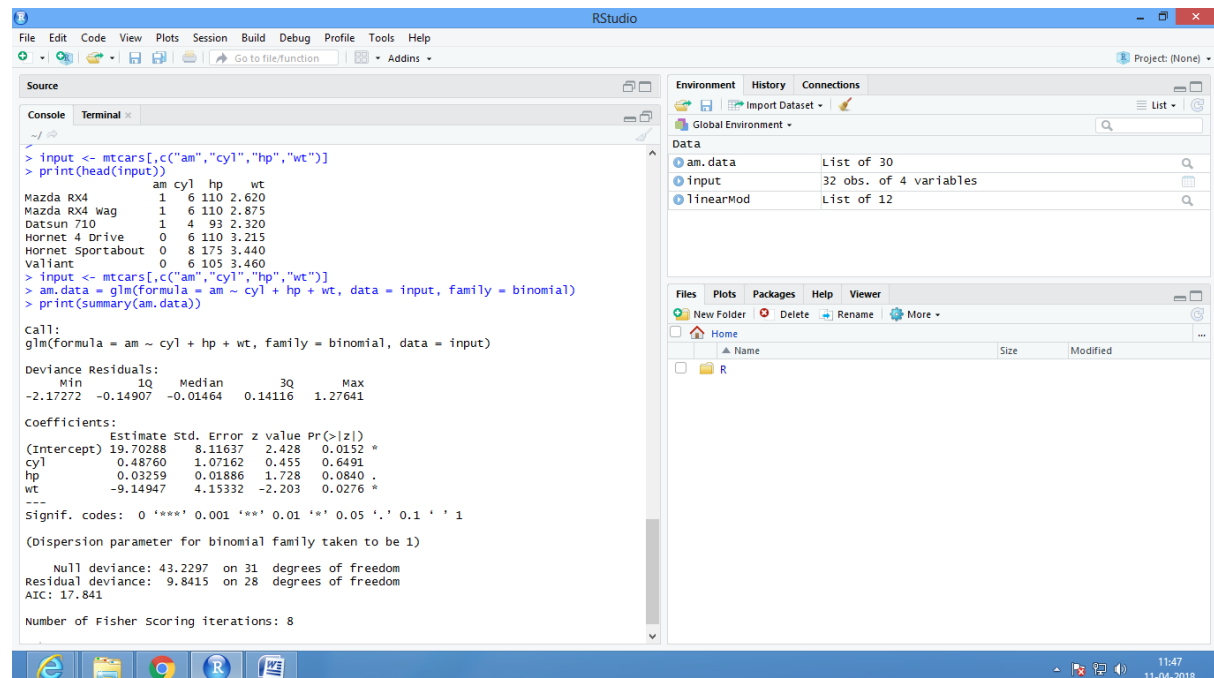
create regression model:

```
input <- mtcars[,c("am","cyl","hp","wt")]
```

```
am.data = glm(formula = am ~ cyl + hp + wt, data = input, family = binomial)
```

```
print(summary(am.data))
```

OUTPUT:



The screenshot displays the RStudio interface with the following components:

- Source Editor:** Contains the R code for selecting columns from the 'mtcars' dataset and fitting a logistic regression model using the 'glm' function.
- Console:** Shows the execution of the code, including the head of the 'input' dataset and the summary of the 'am.data' model.
- Environment:** Lists the objects created in the global environment: 'am.data' (List of 30), 'input' (32 obs. of 4 variables), and 'linearMod' (List of 12).
- Files:** Shows the file explorer with a folder named 'R'.

Console Output:

```
> input <- mtcars[,c("am","cyl","hp","wt")]
> print(head(input))
  am cyl  hp  wt
Mazda RX4      1   6 110 2.620
Mazda RX4 Wag  1   6 110 2.875
Datsun 710      1   4  93 2.320
Hornet 4 Drive  0   6 110 3.215
Hornet Sportabout 0   8 175 3.440
Valiant         0   6 105 3.460
> input <- mtcars[,c("am","cyl","hp","wt")]
> am.data = glm(formula = am ~ cyl + hp + wt, data = input, family = binomial)
> print(summary(am.data))

Call:
glm(formula = am ~ cyl + hp + wt, family = binomial, data = input)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.17272  -0.14907  -0.01464   0.14116   1.27641

Coefficients:
(Intercept)  19.70288   8.11637   2.428   0.0152 *
cyl          0.48760   1.07162   0.455   0.6491
hp           0.03259   0.01886   1.728   0.0840 .
wt          -9.14947   4.15332  -2.203   0.0276 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 43.2297  on 31  degrees of freedom
Residual deviance:  9.8415  on 28  degrees of freedom
AIC: 17.841

Number of Fisher Scoring iterations: 8
```

VIVA QUESTIONS

1. What is linear regression?

Linear regression is a linear approach for modelling the relationship between a scalar response and one or more explanatory variables (also known as dependent and independent variables).

2. What is logistic regression?

Logistic regression is a statistical model that in its basic form uses a logistic function to model a binary dependent variable, although many more complex extensions exist.

3. What are the assumptions of logistic regression?

Basic assumptions that must be met for logistic regression include independence of errors, linearity in the logit for continuous variables, absence of multicollinearity, and lack of strongly influential outliers.

4. Is logistic regression a type of supervised machine learning algorithm?

Logistic regression is a supervised learning classification algorithm used to predict the probability of a target variable.

5. Why logistic regression is called regression and not classification?

Logistic regression uses the same basic formula as linear regression but it is regressing for the probability of a categorical outcome. Linear regression gives a continuous value of output y for a given input X . Whereas, logistic regression gives a continuous value of $P(Y=1)$ for a given input X , which is later converted to $Y=0$ or $Y=1$ based on a threshold value. That's the reason, logistic regression has "Regression" in its name.

WEEK5

IMPLEMENTATION OF SVM CLASSIFICATION TECHNIQUES

IMPLEMENTATION :(SVM)

To use SVM in R, we have a package e1071. The package is not preinstalled, hence one needs to run the line “install.packages(“e1071”) to install the package and then import the package contents using the library command--library(e1071).

R CODE:

```
x=c(1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20)
```

```
y=c(3,4,5,4,8,10,10,11,14,20,23,24,32,34,35,37,42,48,53,60)
```

```
#Create a data frame of the data
```

```
train=data.frame(x,y)
```

```
#Plot the dataset
```

```
plot(train,pch=16)
```

```
#Linear regression
```

```
model <- lm(y ~ x, train)
```

```
#Plot the model using abline
```

```
abline(model)
```

```
#SVM
```

```
library(e1071)
```

```
#Fit a model. The function syntax is very similar to lm function
```

```
model_svm <- svm(y ~ x , train)
```

```
#Use the predictions on the data
```

```
pred <- predict(model_svm, train)
```

```
#Plot the predictions and the plot to see our model fit
```

```
points(train$x, pred, col = "blue", pch=4)
```

```
#Linear model has a residuals part which we can extract and directly calculate rmse
```

```
error <- model$residuals
```

```
lm_error <- sqrt(mean(error^2)) # 3.832974
```

```
#For svm, we have to manually calculate the difference between actual values (train$y) with our predictions (pred)
```

```
error_2 <- train$y - pred
```

```
svm_error <- sqrt(mean(error_2^2)) # 2.696281
```

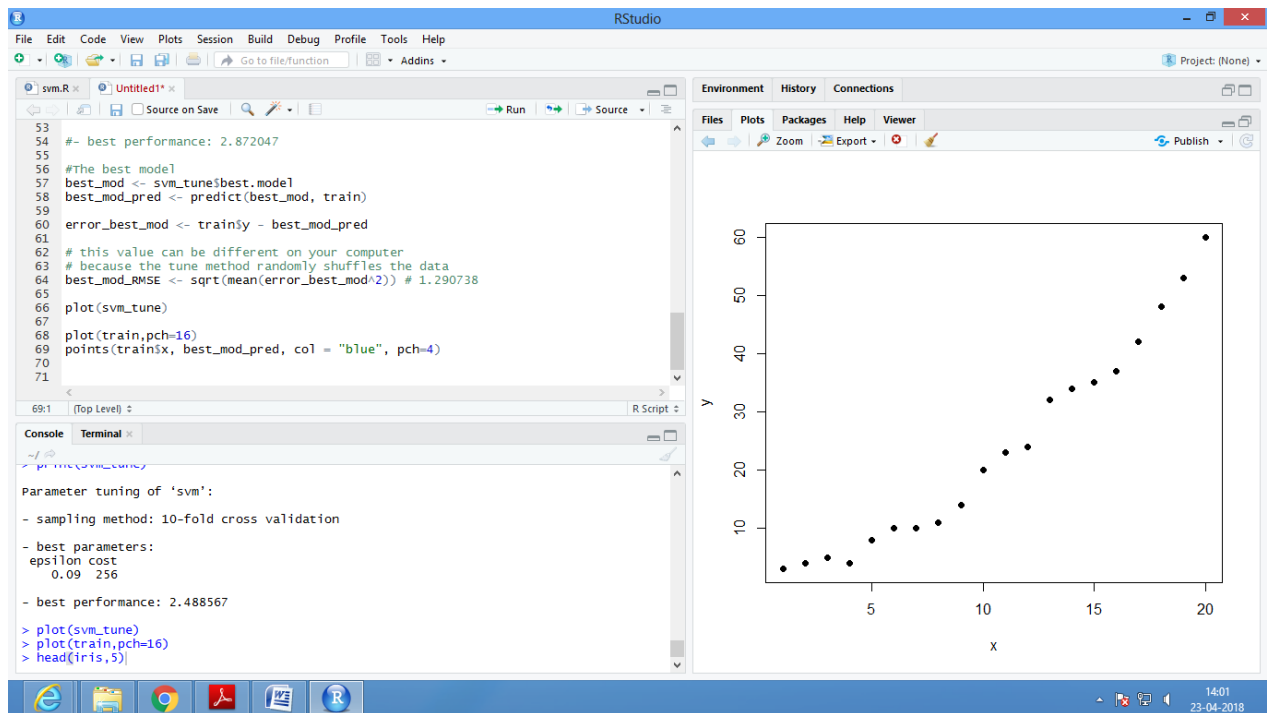
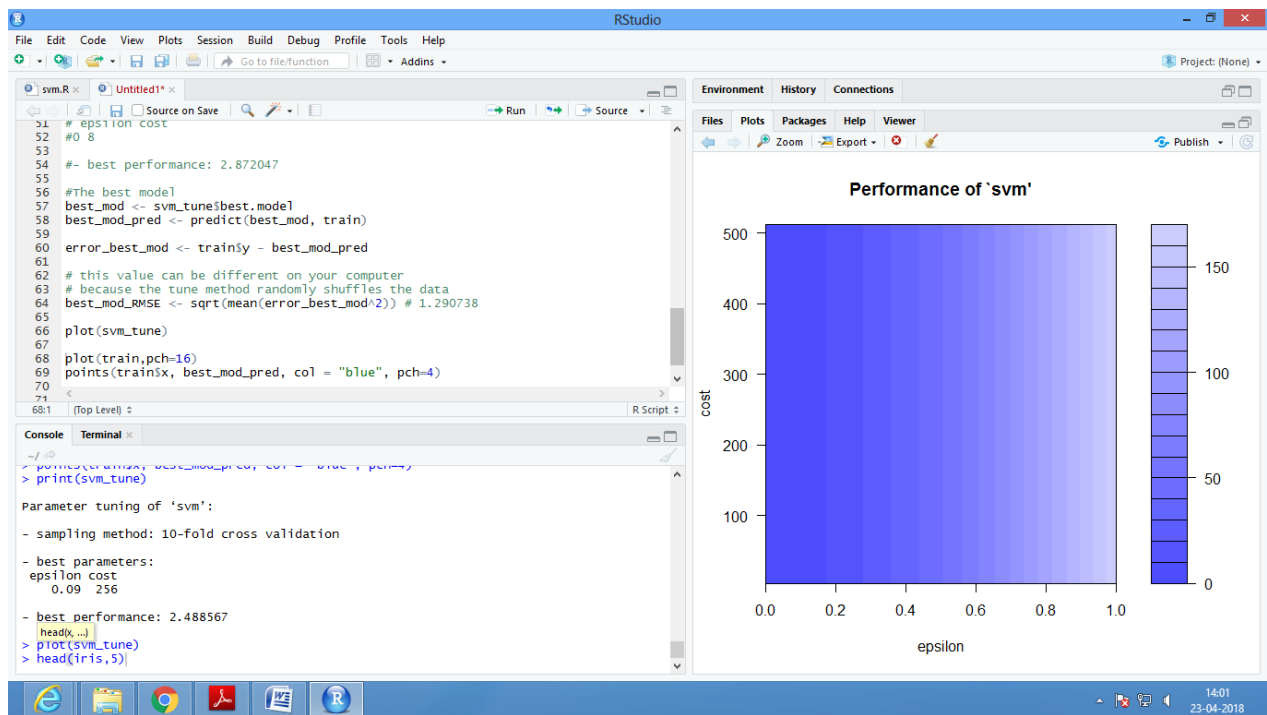
```
# perform a grid search
```

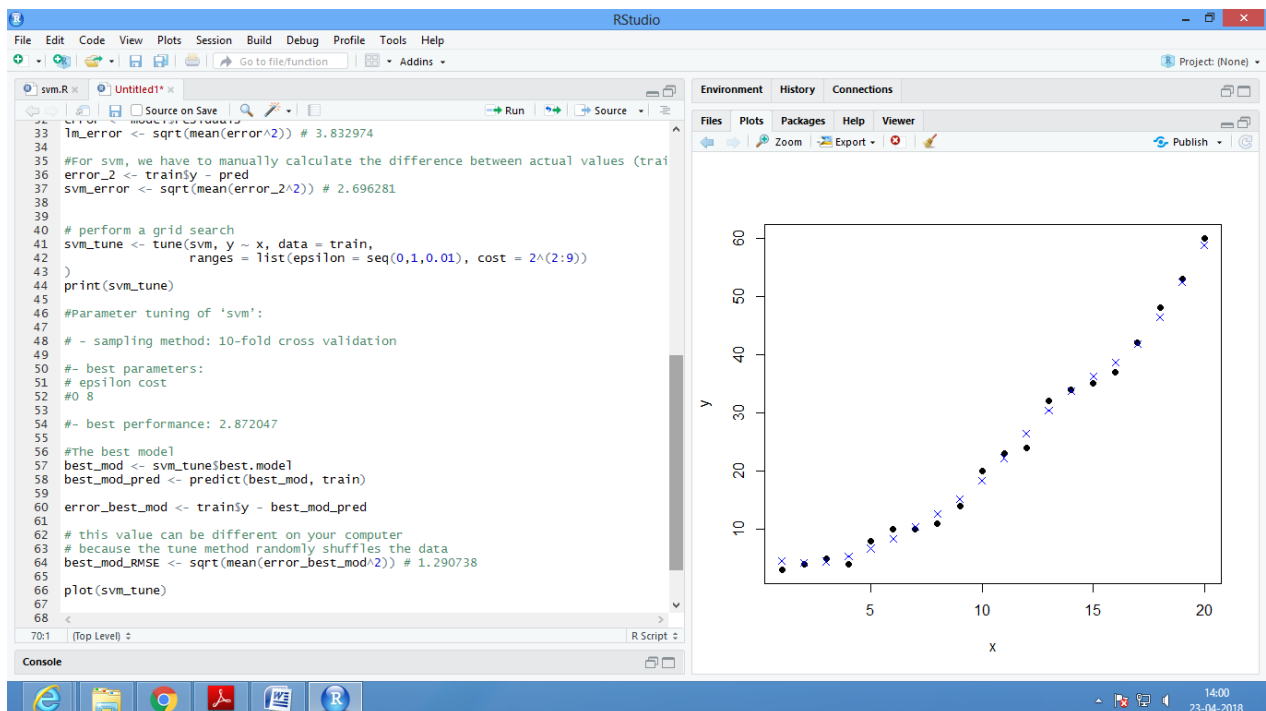
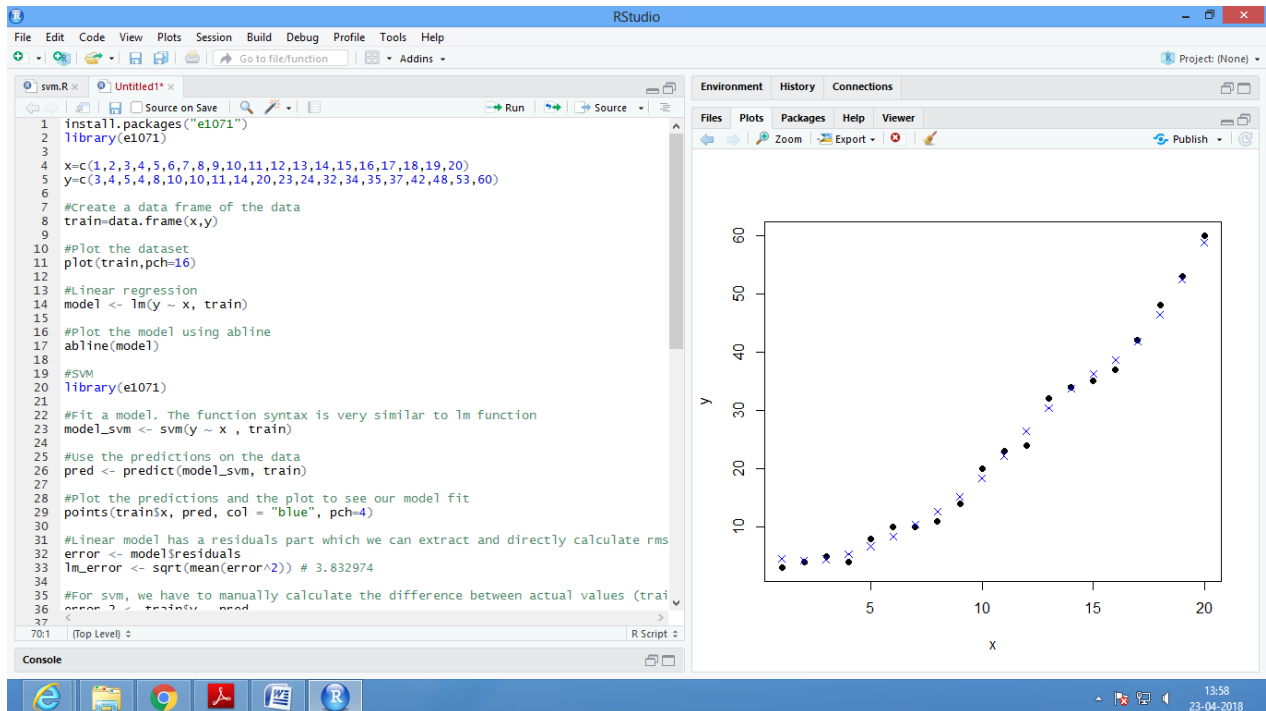
```
svm_tune <- tune(svm, y ~ x, data = train,
```

```
ranges = list(epsilon = seq(0,1,0.01), cost = 2^(2:9))
```

```
)  
print(svm_tune)  
  
#Parameter tuning of 'svm':  
  
# - sampling method: 10-fold cross validation  
  
#- best parameters:  
# epsilon cost  
#0 8  
  
#- best performance: 2.872047  
  
#The best model  
best_mod <- svm_tune$best.model  
best_mod_pred <- predict(best_mod, train)  
  
error_best_mod <- train$y - best_mod_pred  
  
# this value can be different on your computer  
# because the tune method randomly shuffles the data  
best_mod_RMSE <- sqrt(mean(error_best_mod^2)) # 1.290738  
  
plot(svm_tune)  
  
plot(train,pch=16)  
points(train$x, best_mod_pred, col = "blue", pch=4)
```

OUTPUT :





VIVA QUESTIONS

1. How SVM is used for classification?

SVM is a supervised machine learning algorithm which can be used for classification or regression problems. It uses a technique called the kernel trick to transform your data and then based on these transformations it finds an optimal boundary between the possible outputs.

2. Why is SVM used?

SVM can be used for classification (distinguishing between several groups or classes) and regression (obtaining a mathematical model to predict something). They can be applied to both linear and nonlinear problems.

3. What is the fundamental idea of SVM for classification?

SVM or Support Vector Machine is a linear model for classification and regression problems. It can solve linear and non-linear problems and work well for many practical problems.

4. Give some real-world applications of SVM?

Inverse Geosounding Problem, Texture Classification, Speech Recognition and more..

5. Explain about SVM regression?

Support Vector Regression is a supervised learning algorithm that is used to predict discrete values. Support Vector Regression uses the same principle as the SVMs. The basic idea behind SVR is to find the best fit line.

WEEK6

IMPLEMENTATION OF DECISION TREE CLASSIFICATION TECHNIQUES

Use the below command in R console to install the package. You also have to install the dependent packages if any.

install.packages("party")

The basic syntax for creating a decision tree in R is

ctree(formula, data)

input data:

We will use the R in-built data set named readingSkills to create a decision tree. It describes the score of someone's readingSkills if we know the variables "age","shoesize","score" and whether the person is a native speaker or not.

Load the party package. It will automatically load other dependent packages.

library(party)

Print some records from data set readingSkills.

print(head(readingSkills))

Load the party package. It will automatically load other dependent packages.

We will use the ctree() function to create the decision tree and see its graph.

Load the party package. It will automatically load other dependent packages.

library(party)

Create the input data frame.

input.dat <- readingSkills[c(1:105),]

Give the chart file a name.

png(file = "decision_tree.png")

Create the tree.

output.tree <- ctree(nativeSpeaker ~ age + shoeSize + score, data = input.dat)

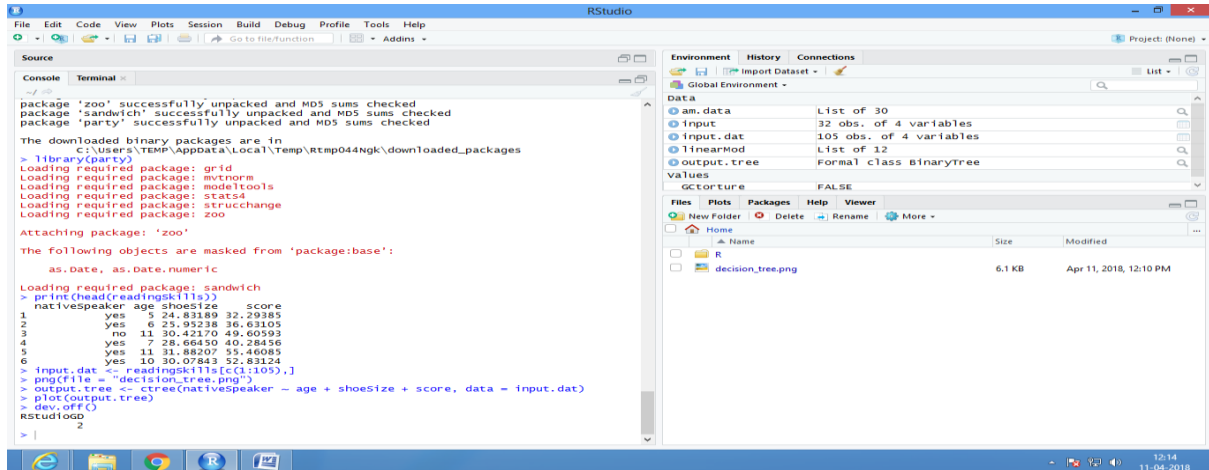
Plot the tree.

plot(output.tree)

Save the file.

dev.off()

OUTPUT:



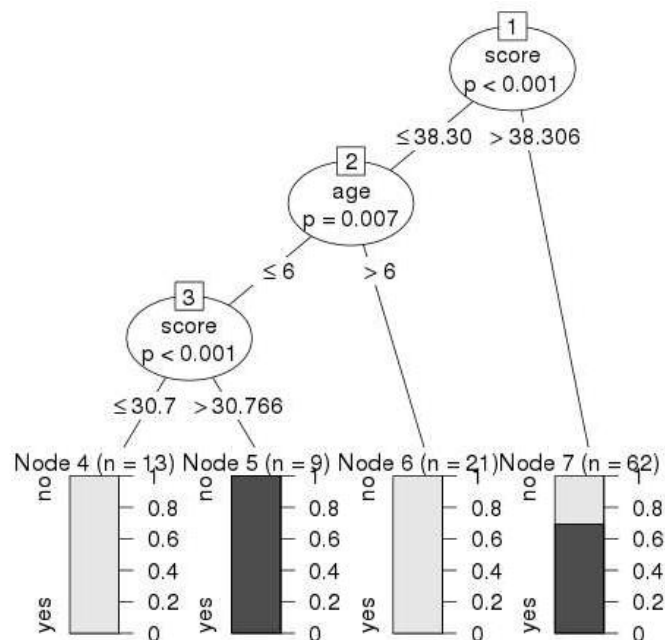
```
package 'zoo' successfully unpacked and MD5 sums checked
package 'sandwich' successfully unpacked and MD5 sums checked
package 'party' successfully unpacked and MD5 sums checked
The downloaded binary packages are in
C:\Users\R\TEMP\AppData\Local\Temp\044ngk\downloaded_packages
> library(party)
Loading required package: grid
Loading required package: mvtnorm
Loading required package: modeltools
Loading required package: stats4
Loading required package: strucchange
Loading required package: zoo
Attaching package: 'zoo'
The following objects are masked from 'package:base':
  as.Date, as.Date.numeric
Loading required package: sandwich
> print(head(readingskills))
  nativespeaker age shoeSize score
1         yes   5 24.83189 32.29385
2         yes   6 25.95238 36.63105
3         no   11 30.42170 49.60593
4         yes   7 28.66450 40.28456
5         yes  11 31.88207 55.46085
6         yes  10 30.07843 52.83124
> input.dat <- readingskills[c(1:105),]
> png(file = "decision_tree.png")
> output.tree <- ctree(nativespeaker ~ age + shoeSize + score, data = input.dat)
> plot(output.tree)
> dev.off()
RStudio 2
```

Environment

| Object | Class |
|-------------|-------------------------|
| am.data | List of 30 |
| input | 32 obs. of 4 variables |
| input.dat | 105 obs. of 4 variables |
| linearMod | List of 12 |
| output.tree | Formal class ctree |

Files

| Name | Size | Modified |
|-------------------|--------|------------------------|
| Home | | |
| R | | |
| decision_tree.png | 6.1 KB | Apr 11, 2018, 12:10 PM |



VIVA QUESTIONS

1. Why are decisions trees the most popular classification technique?

Decision trees are able to generate understandable rules. Decision trees perform classification without requiring much computation. Decision trees are able to handle both continuous and categorical variables. Decision trees provide a clear indication of which fields are most important for prediction or classification.

2. Is decision tree a classification problem?

Decision trees are a rule-based approach to classification and regression problems. They use the values in each feature to split the dataset to a point where all data points that have the same class are grouped together.

3. Which algorithm is used in decision tree?

The basic algorithm used in decision trees is known as the ID3 (by Quinlan) algorithm. The ID3 algorithm builds decision trees using a top-down, greedy approach.

4. How decision trees used for classification?

Decision tree builds classification or regression models in the form of a tree structure. It breaks down a dataset into smaller and smaller subsets while at the same time an associated decision tree is incrementally developed.

5. What is information gain?

Information gain is the reduction in entropy or surprise by transforming a dataset and is often used in training decision trees. Information gain is calculated by comparing the entropy of the dataset before and after a transformation.

WEEK7

IMPLEMENTATION OF HIERARCHICAL CLUSTERING

PROGRAM:

Installing and loading required R packages

```
install.packages("factoextra")
```

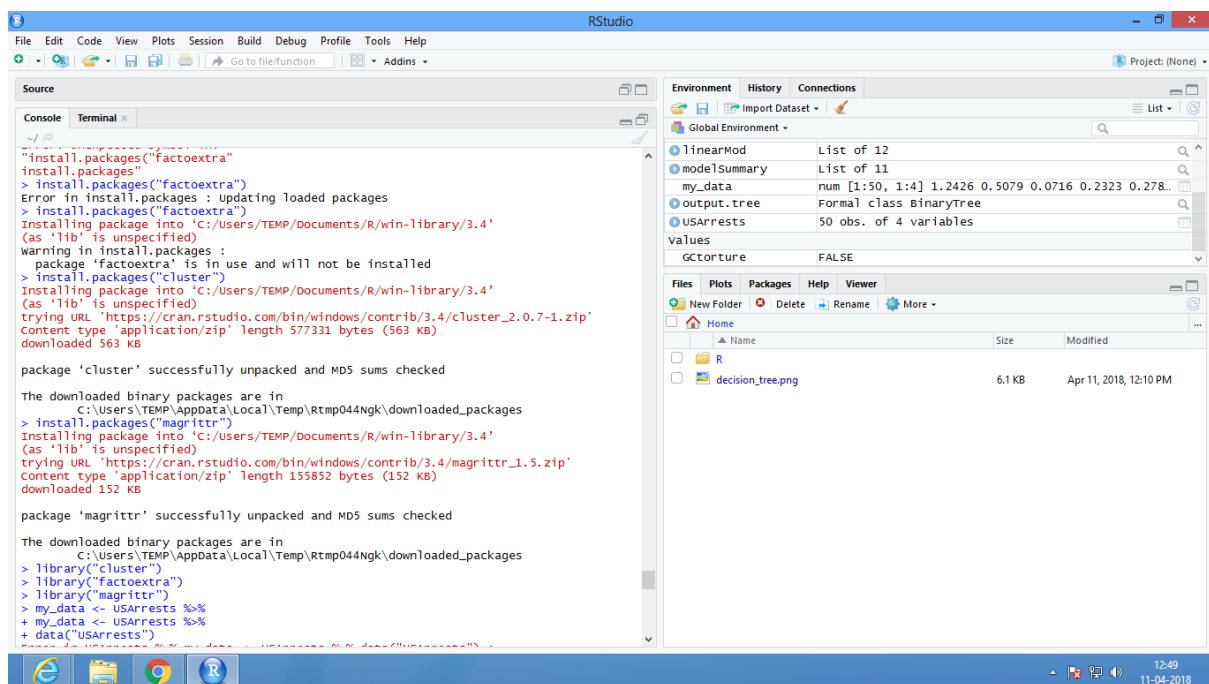
```
install.packages("cluster")
```

```
install.packages("magrittr")
```

```
library("cluster")
```

```
library("factoextra")
```

```
library("magrittr")
```



Data preparation

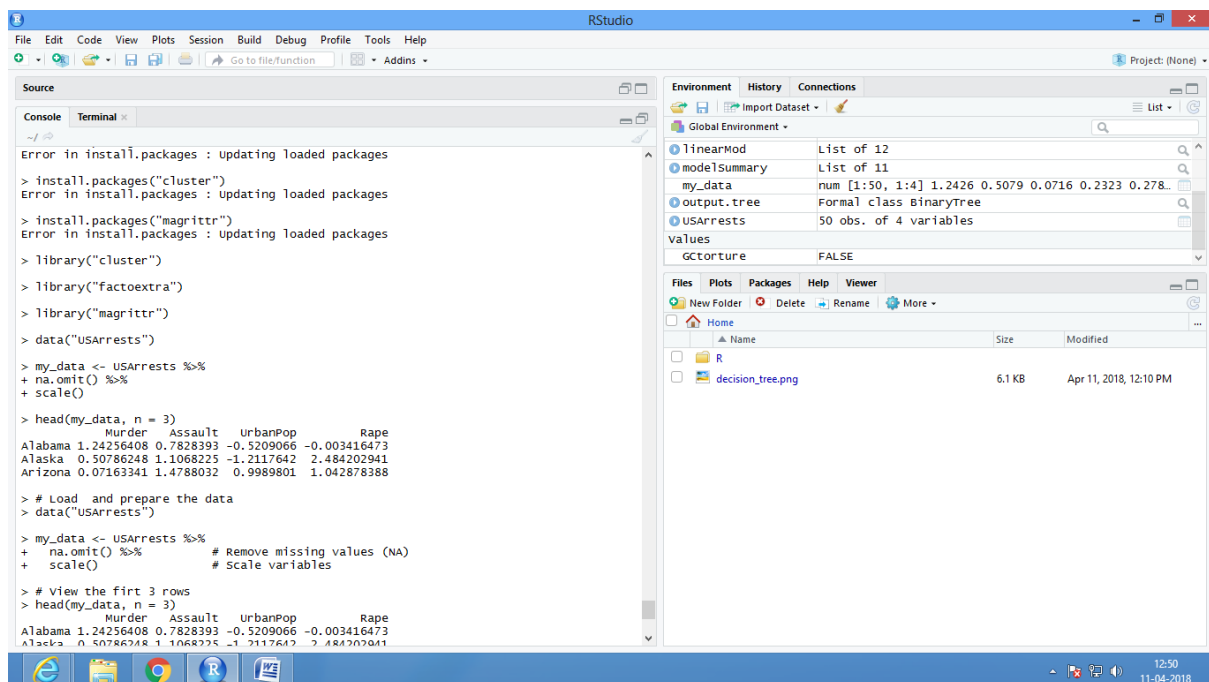
```
# Load and prepare the data
```

```
data("USArrests")
```

```
my_data <- USArrests %>% na.omit() %>% # Remove missing values (NA) scale() #  
Scale variables
```

View the first 3 rows

head(my_data, n = 3)



The screenshot shows the RStudio interface. The console on the left displays the following code and output:

```
Error in install.packages : updating loaded packages
> install.packages("cluster")
Error in install.packages : updating loaded packages
> install.packages("magrittr")
Error in install.packages : updating loaded packages
> library("cluster")
> library("factoextra")
> library("magrittr")
> data("USArrests")
> my_data <- USArrests %>%
+   na.omit() %>%
+   scale()
> head(my_data, n = 3)
      Murder  Assault  UrbanPop   Rape
Alabama 1.24256408 0.7828393 -0.5209066 -0.003416473
Alaska  0.50786248 1.1068225 -1.2117642  2.484202941
Arizona 0.07163341 1.4788032  0.9989801  1.042878388
> # Load and prepare the data
> data("USArrests")
> my_data <- USArrests %>%
+   na.omit() %>%      # Remove missing values (NA)
+   scale()            # Scale variables
> # View the first 3 rows
> head(my_data, n = 3)
      Murder  Assault  UrbanPop   Rape
Alabama 1.24256408 0.7828393 -0.5209066 -0.003416473
Alaska  0.50786248 1.1068225 -1.2117642  2.484202941
Arizona 0.07163341 1.4788032  0.9989801  1.042878388
```

The Environment pane on the right shows the following objects:

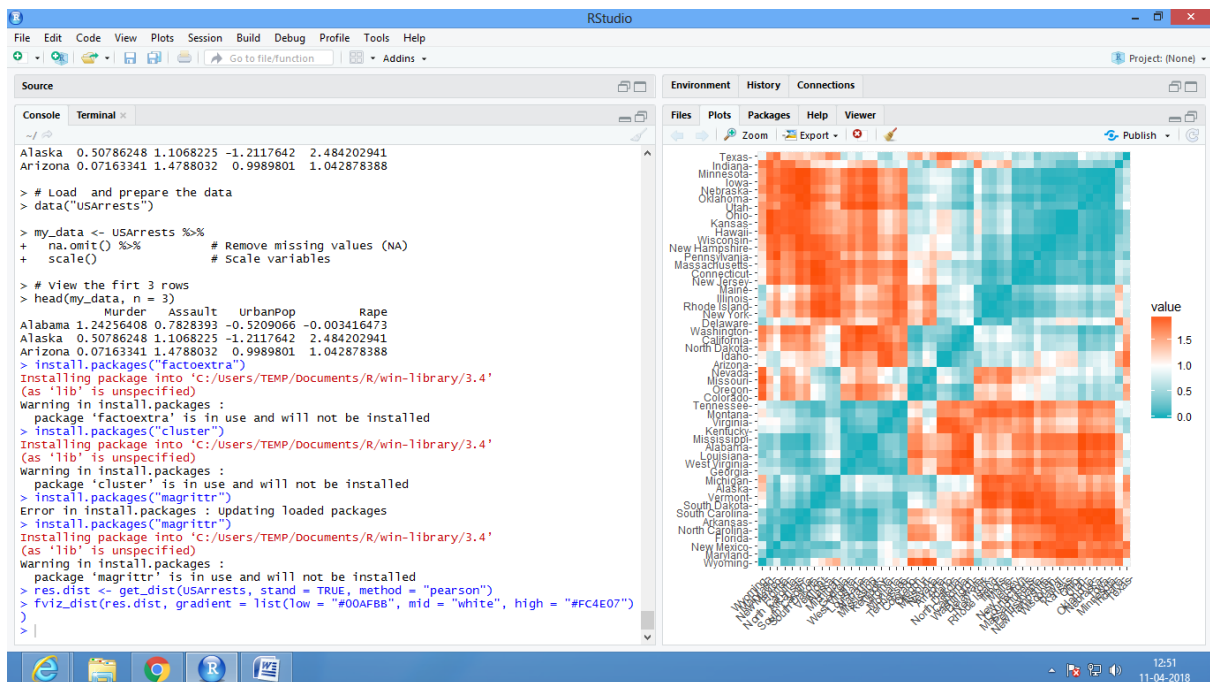
| Object | Type |
|--------------|--|
| LinearMod | List of 12 |
| modelSummary | List of 11 |
| my_data | num [1:50, 1:4] 1.2426 0.5079 0.0716 0.2323 0.278... |
| output.tree | Formal class Binarytree |
| USArrests | 50 obs. of 4 variables |
| values | |
| Gctorture | FALSE |

The Files pane shows the following files:

| Name | Size | Modified |
|-------------------|--------|------------------------|
| R | | |
| decision_tree.png | 6.1 KB | Apr 11, 2018, 12:10 PM |

Distance measures

```
res.dist <- get_dist(USArrests, stand = TRUE, method = "pearson")
fviz_dist(res.dist, gradient = list(low = "#00AFBB", mid = "white", high = "#FC4E07"))
```



VIVA QUESTIONS

1. What is Hierarchical Clustering algorithm?

HCA is an unsupervised clustering algorithm which involves creating clusters that have predominant ordering from top to bottom.

2. What are the various types of Hierarchical clustering algorithm?

Divisive and Agglomerative.

3. What do you mean by clustering?

Clustering is the task of dividing the population or data points into a number of groups such that data points in the same groups are more similar to other data points in the same group and dissimilar to the data points in other groups.

4. What is the minimum number of variables/features required to perform clustering?

At least a single variable is required to perform clustering analysis. Clustering analysis with a single variable can be visualized with the help of a histogram.

5. What is the strength and weaknesses of Hierarchical Clustering?

Strength: sums up the data, good for small data sets.

Weakness: computationally demanding, fails on larger sets.

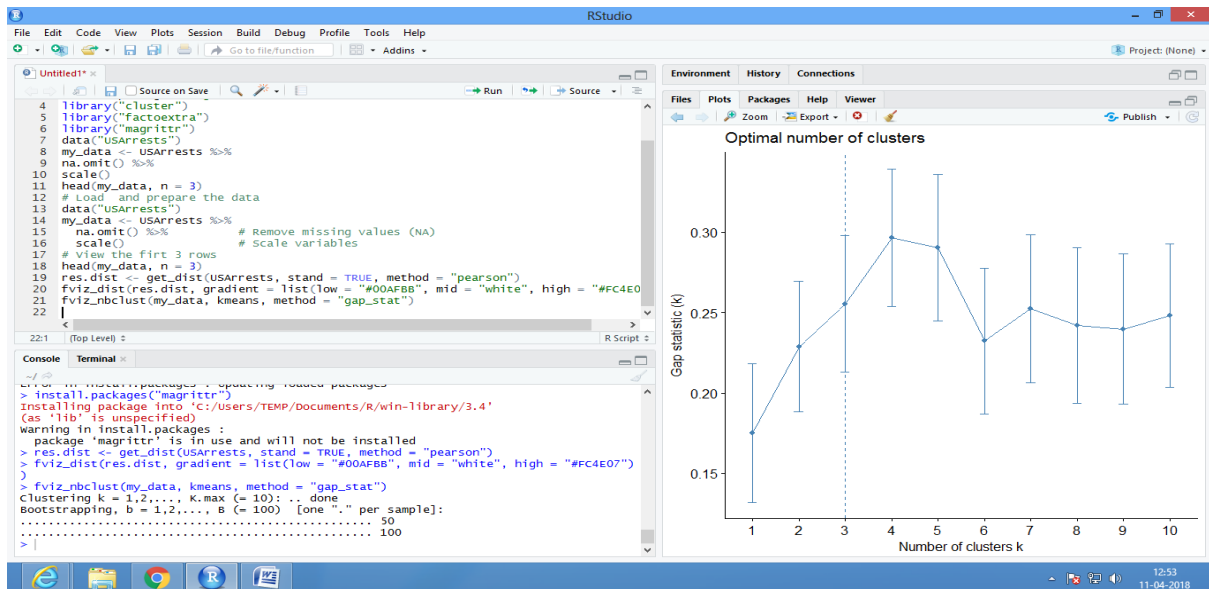
WEEK 8

IMPLEMENTATION OF PARTITIONING CLUSTERING

Determining the optimal number of clusters: use factoextra::fviz_nbclust()

```
library("factoextra")
```

```
fviz_nbclust(my_data, kmeans, method = "gap_stat")
```



Compute and visualize k-means clustering

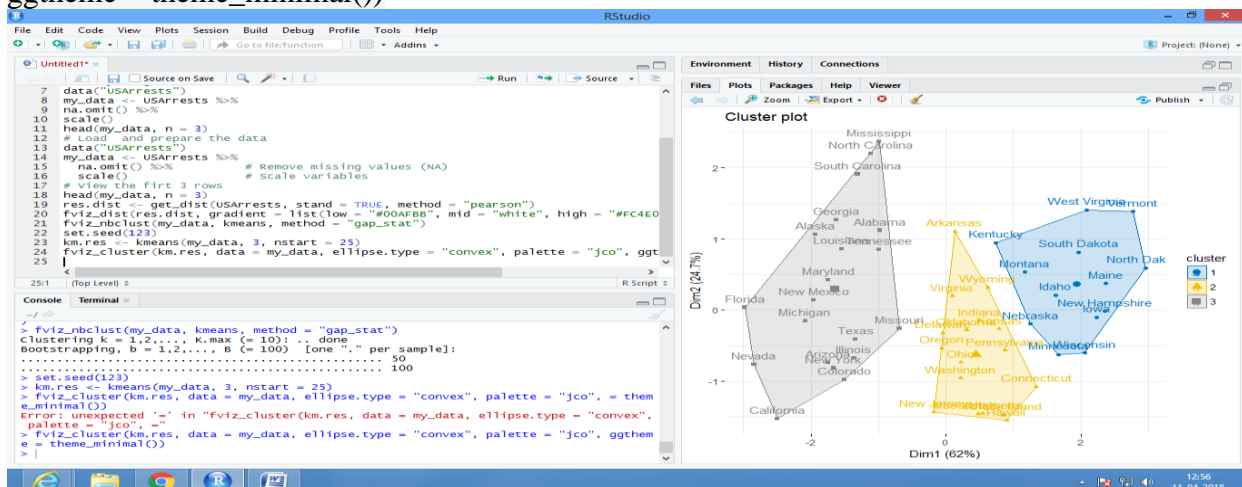
```
set.seed(123)
```

```
km.res <- kmeans(my_data, 3, nstart = 25)
```

```
# Visualize
```

```
library("factoextra")
```

```
fviz_cluster(km.res, data = my_data, ellipse.type = "convex", palette = "jco",  
ggtheme = theme_minimal())
```



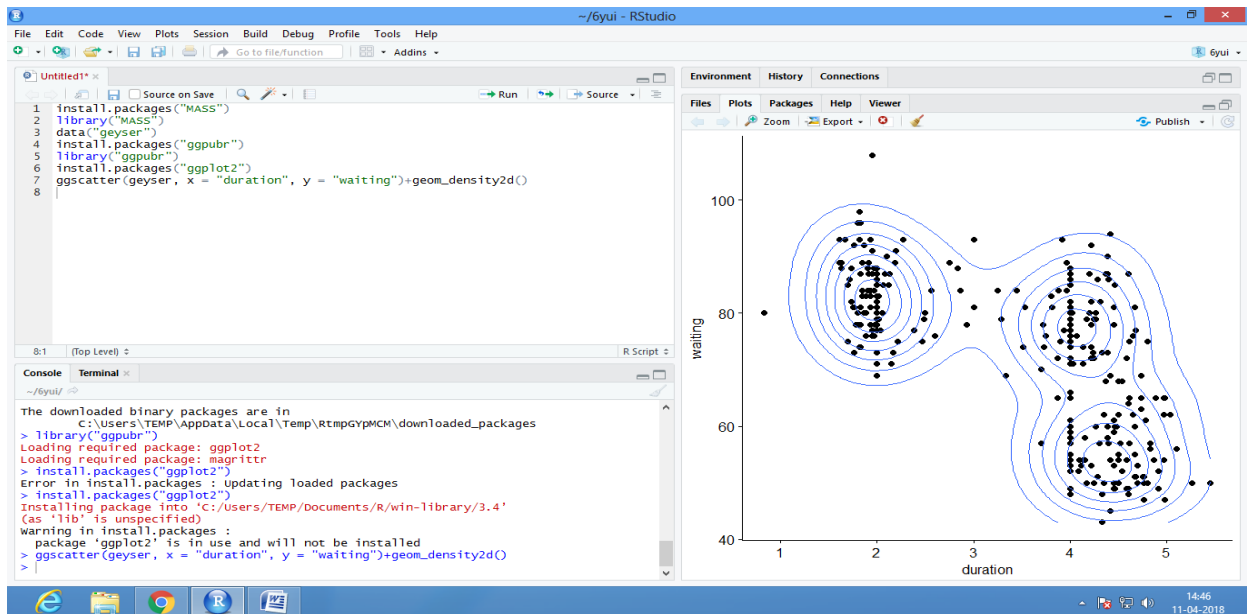
MODEL BASED CLUSTERING:

Load the data

```
library("MASS") data("geyser")
```

Scatter plot **library("ggpubr")**

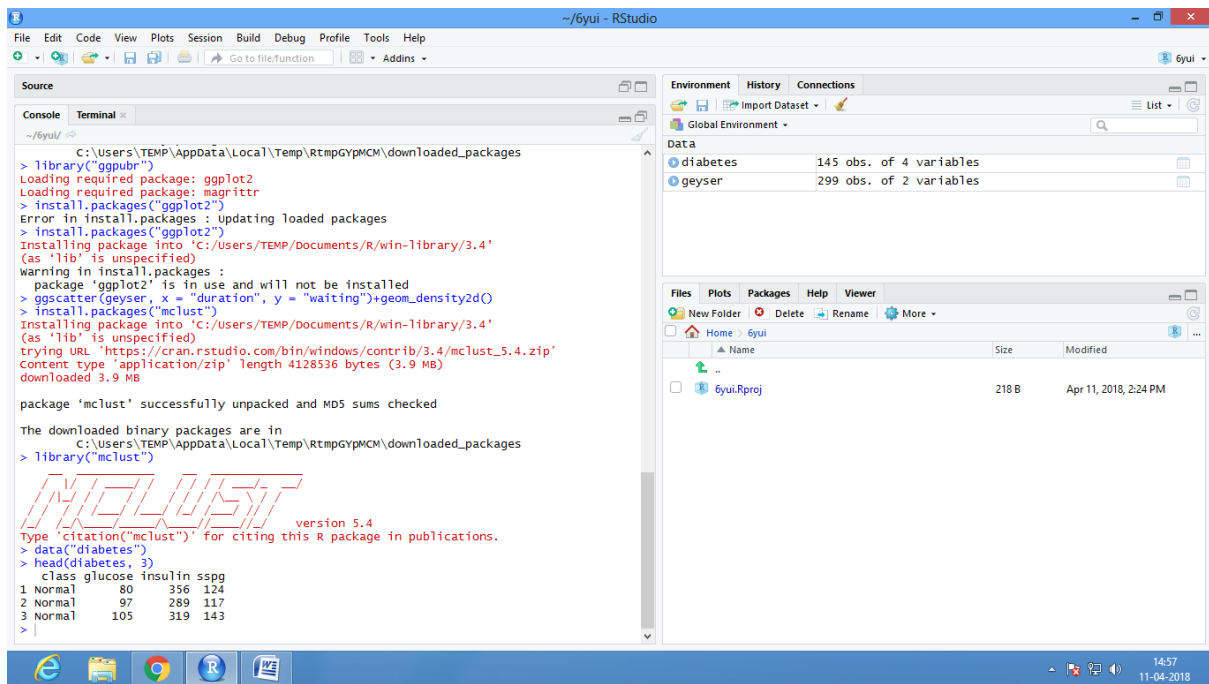
```
ggscatter(geyser, x = "duration", y = "waiting")+ geom_density2d() # Add 2D density
```



```
library("mclust")
```

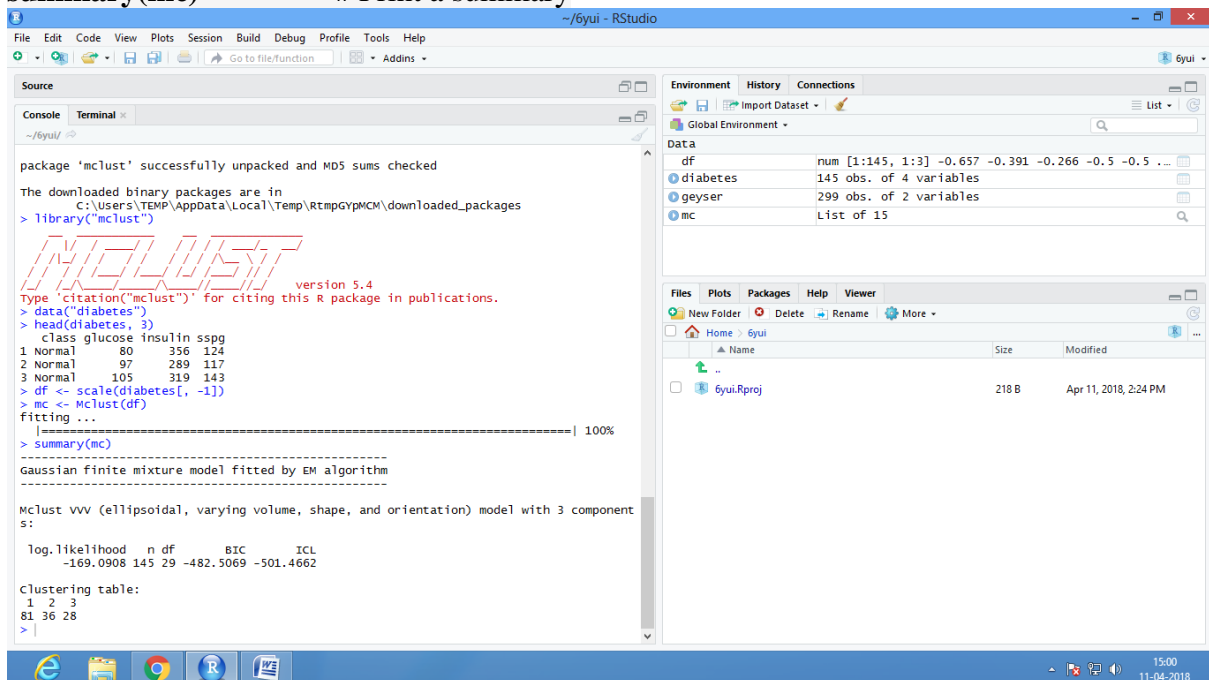
```
data("diabetes")
```

```
head(diabetes, 3)
```



Model-based clustering can be computed using the function `Mclust()` as follow:

```
library(mclust)
df <- scale(diabetes[, -1]) # Standardize the data
mc <- Mclust(df)           # Model-based-clustering
summary(mc)                # Print a summary
```



mc\$modelName # Optimal selected model ==> "VVV"

mc\$G # Optimal number of cluster => 3 **head(mc\$z, 30)** # Probability to belong to a given cluster **head(mc\$classification, 30)** # Cluster assignement of each observation

The screenshot shows the RStudio interface with the following components:

- Source:** Contains the R code being executed:

```
> mc$G
[1] 3
> head(mc$z, 30)
      [,1]      [,2]      [,3]
1 0.9906745 0.008991332 3.341728e-04
2 0.9822128 0.017783229 3.974744e-06
3 0.9777871 0.022157665 5.527579e-05
4 0.9774763 0.022312280 2.113743e-04
5 0.9208978 0.079034264 6.789759e-05
6 0.9863472 0.012977950 6.748263e-04
7 0.9429761 0.056900375 1.235066e-04
8 0.9768888 0.022939764 1.714394e-04
9 0.9739547 0.025861249 1.840708e-04
10 0.9841312 0.015858549 1.028654e-05
11 0.9603523 0.039514294 1.334497e-04
12 0.9818178 0.018066502 1.156881e-04
13 0.9815268 0.018323709 1.494696e-04
14 0.9748375 0.024856829 3.056499e-04
15 0.9672956 0.032528109 1.762948e-04
16 0.9755170 0.020741317 3.741703e-03
17 0.9865950 0.012222799 1.821523e-04
18 0.9716530 0.028157524 1.894882e-04
19 0.9852762 0.014713123 1.068929e-05
20 0.9877039 0.012263455 3.266338e-05
21 0.9842866 0.015504504 2.089191e-04
22 0.9872129 0.012765078 2.204329e-05
23 0.9797791 0.020088555 1.323722e-04
24 0.9801459 0.019729065 1.250125e-04
25 0.6561084 0.343880540 1.110662e-05
26 0.9550401 0.039404741 5.555189e-03
27 0.9778729 0.022010199 1.168983e-04
28 0.9759071 0.023941785 1.511203e-04
29 0.9847403 0.015082453 1.772440e-04
30 0.9494749 0.050440377 8.472274e-05
> head(mc$classification, 30)
 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28
1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1
```
- Environment:** Shows the global environment with variables: **df** (num [1:145, 1:3] -0.657 -0.391 -0.266 -0.5 -0.5 ...), **diabetes** (145 obs. of 4 variables), **geyser** (299 obs. of 2 variables), and **mc** (List of 15).
- Files:** Shows the file explorer with a folder named **6yui.Rproj** (218 B, Apr 11, 2018, 2:24 PM).

VISUALIZING MODEL-BASED CLUSTERING

library(factoextra)

BIC values used for choosing the number of clusters

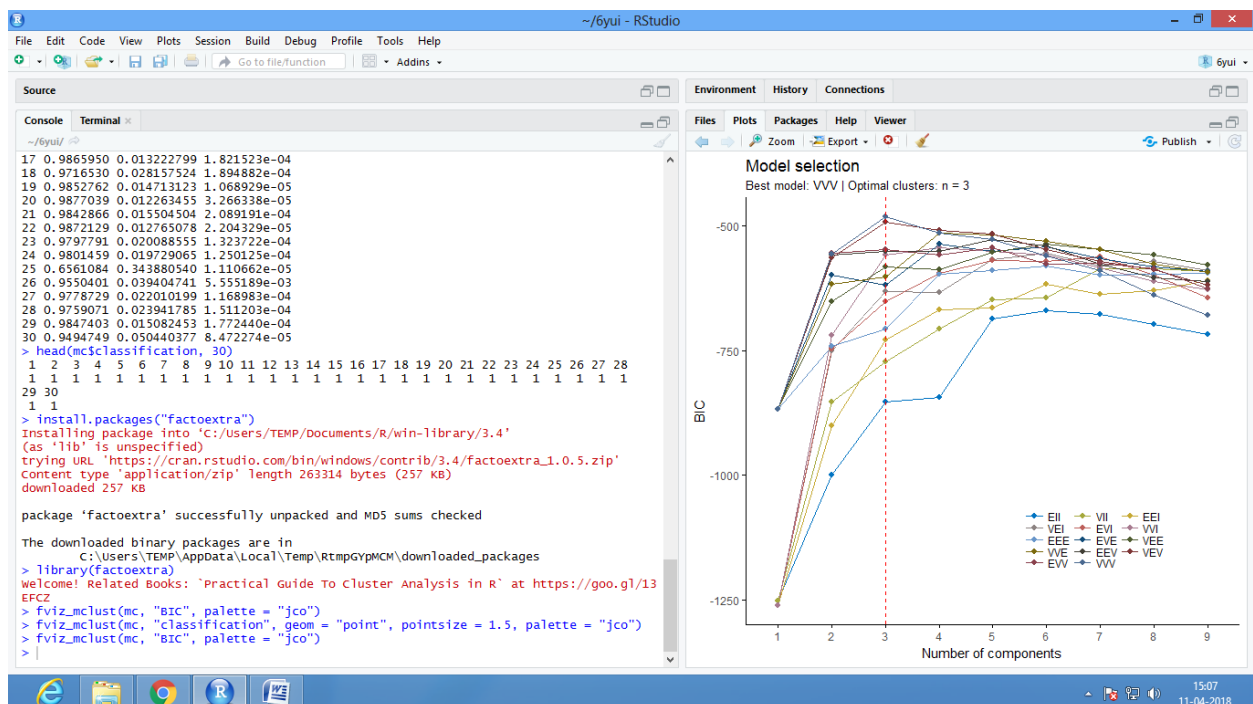
```
fviz_mclust(mc, "BIC", palette = "jco")
```

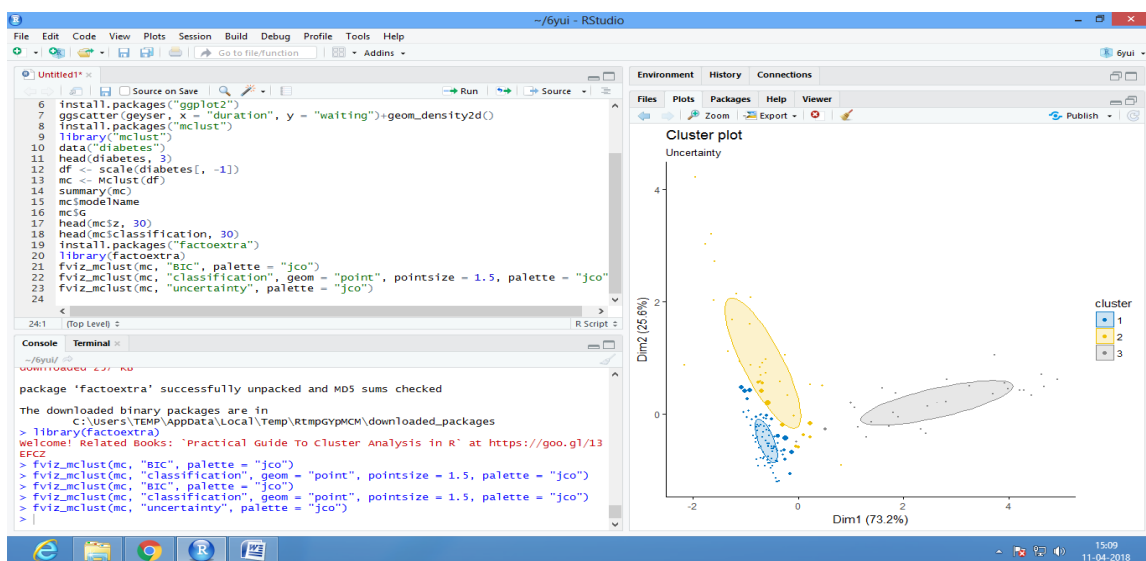
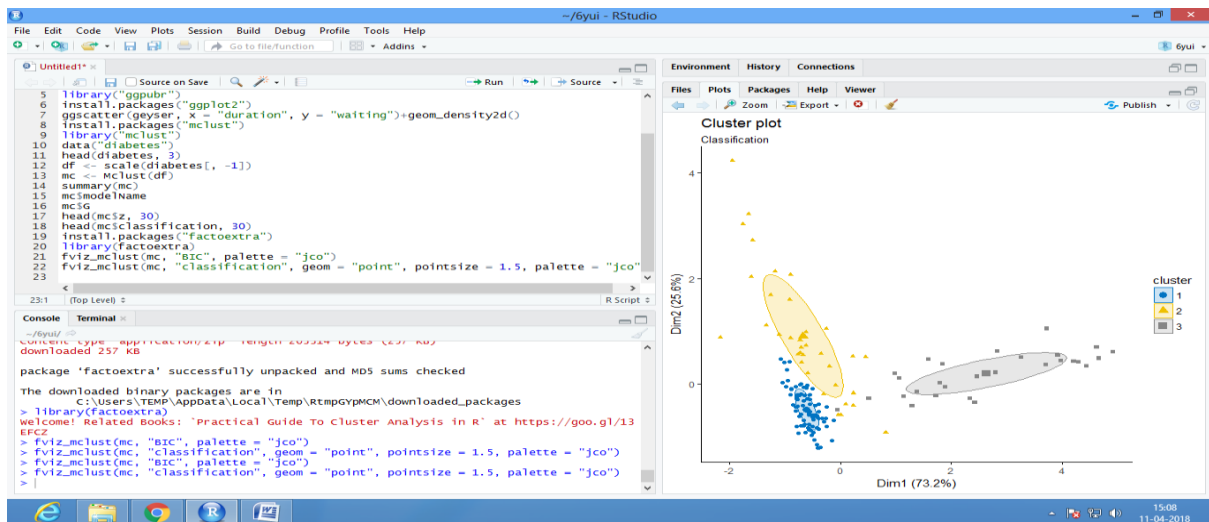
Classification: plot showing the clustering

```
fviz_mclust(mc, "classification", geom = "point", pointsize = 1.5, palette = "jco")
```

Classification uncertainty

```
fviz_mclust(mc, "uncertainty", palette = "jco")
```





VIVA QUESTIONS

1. What is partitioning in clustering?

Partitional clustering (or partitioning clustering) are clustering methods used to classify observations, within a data set, into multiple groups based on their similarity. The algorithms require the analyst to specify the number of clusters to be generated.

2. What is the purpose of K-Mean Clustering?

The K-means clustering algorithm is used to find groups which have not been explicitly labeled in the data. This can be used to confirm business assumptions about what types of groups exist or to identify unknown groups in complex data sets.

3. What is clustering used for?

Clustering is an unsupervised machine learning method of identifying and grouping similar data points in larger datasets without concern for the specific outcome.

4. How does K-mean Clustering work?

K-means clustering uses “centroids”, K different randomly-initiated points in the data, and assigns every data point to the nearest centroid.

5. What are the advantages and disadvantages of k-means algorithm?

Advantages

- 1) If variables are huge, then K-Means most of the times computationally faster than hierarchical clustering, if we keep k smalls.
- 2) K-Means produce tighter clusters than hierarchical clustering, especially if the clusters are globular.

Disadvantages

- 1) Difficult to predict K-Value.
- 2) With global cluster, it didn't work well.
- 3) Different initial partitions can result in different final clusters.
- 4) It does not work well with clusters (in the original data) of Different size and Different density

WEEK9

DATA VISUALIZATION USING PIE, BAR, BOXPLOT CHART PLOTTING FRAMEWORK

IMPLEMENTATION:

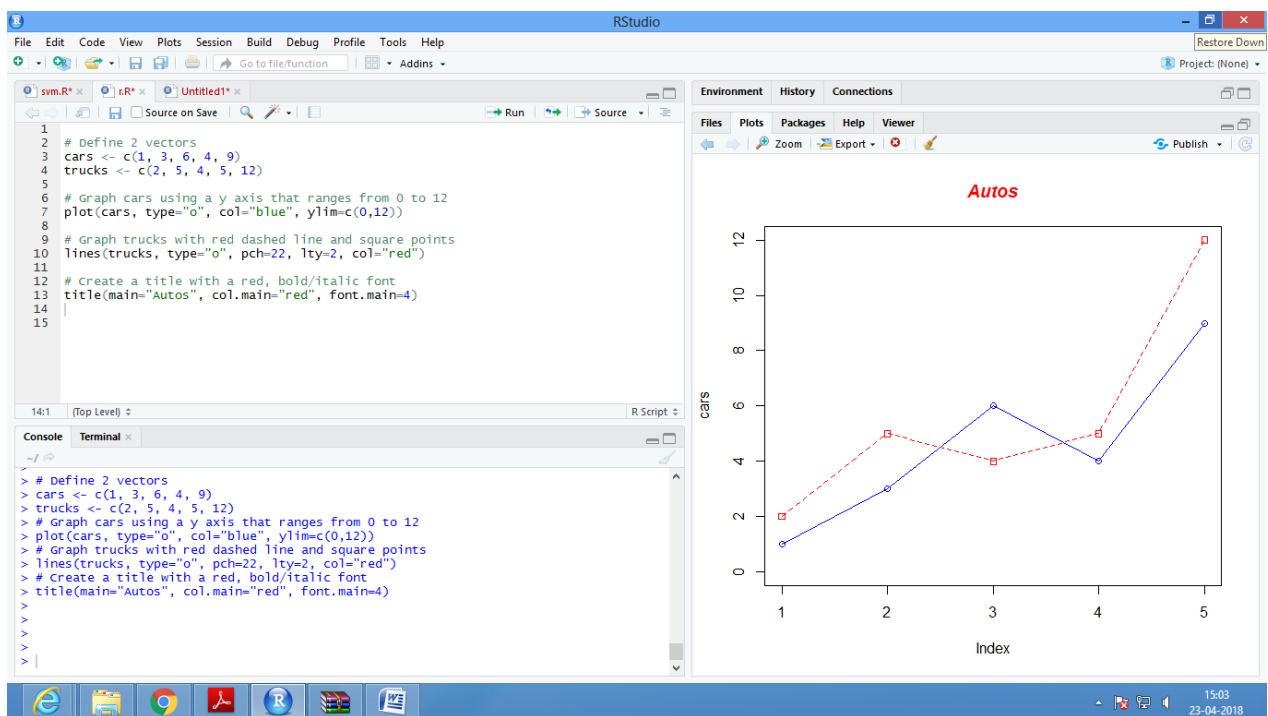
Step 1: Define two vectors as truck and car

Step 2: Using plot function specify the y-axis range directly so it will be large enough to fit the truck data

Step 3: plot cars using a y axis that ranges from 0 to 12

Step 4: plot trucks with red dashed line and square point trucks with red dashed line and square points

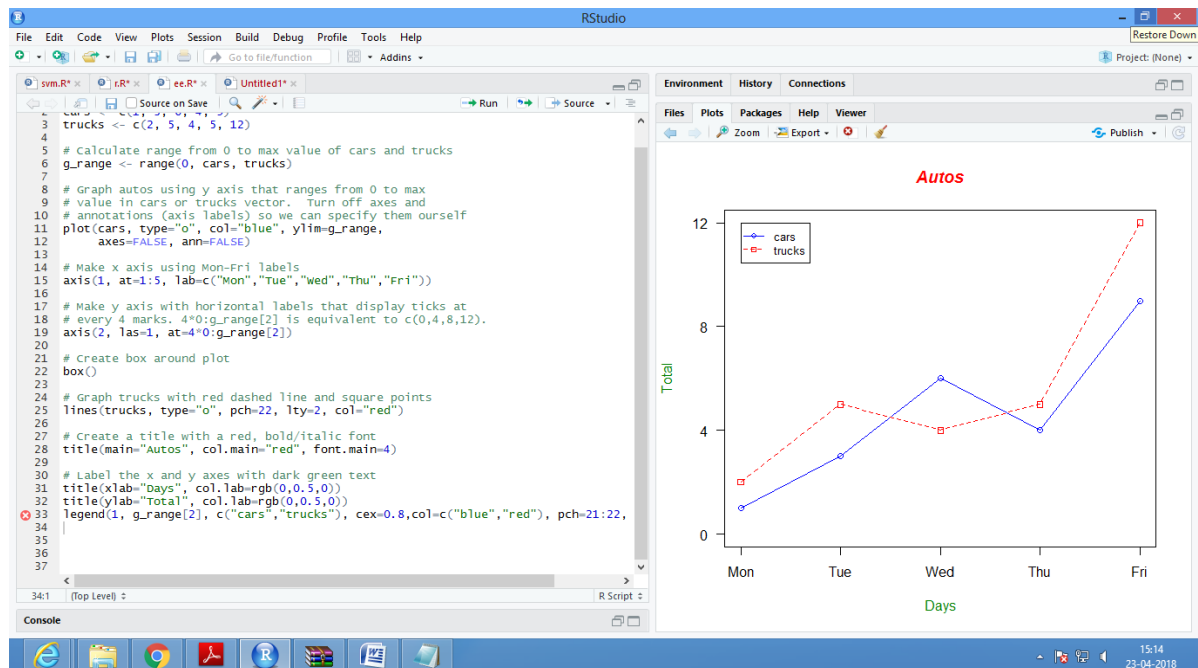
Step 5: Create a title with a red, bold/italic font



Step 6: Create a legend at (1, g_range[2]) that is slightly smaller (cex) and uses the same line colors and points used by the actual plots.


```
legend(1, g_range[2], c("cars","trucks"), cex=0.8,col=c("blue","red"), pch=21:22, lty=1:2);
```

X`



VIVA QUESTIONS

1. What is data visualization?

The representation of information in the form of a chart, diagram, picture, etc.

2. What are some important features of good data visualization?

- It's visually appealing. ...
- It's scalable. ...
- It gives the user the right information. ...
- It's accessible. ...
- It allows rapid development and deployment.

3. What is a bar plot? For what type of data is bar plot usually used for?

Bar charts are one of the many techniques used to present data in a visual form so that the reader may readily recognize patterns or trends. Bar charts usually present categorical variables, discrete variables or continuous variables grouped in class intervals.

4. What is an outlier?

Outlier is an observation of data that does not fit the rest of the data. It is sometimes called an extreme value. When you graph an outlier, it will appear not to fit the pattern of the graph.

5. What type of data is box-plots usually used for? Why?

Box plots are used to show distributions of numeric data values, especially when you want to compare them between multiple groups. They are built to provide high-level information at a glance, offering general information about a group of data's symmetry, skew, variance, and outliers.

WEEK10

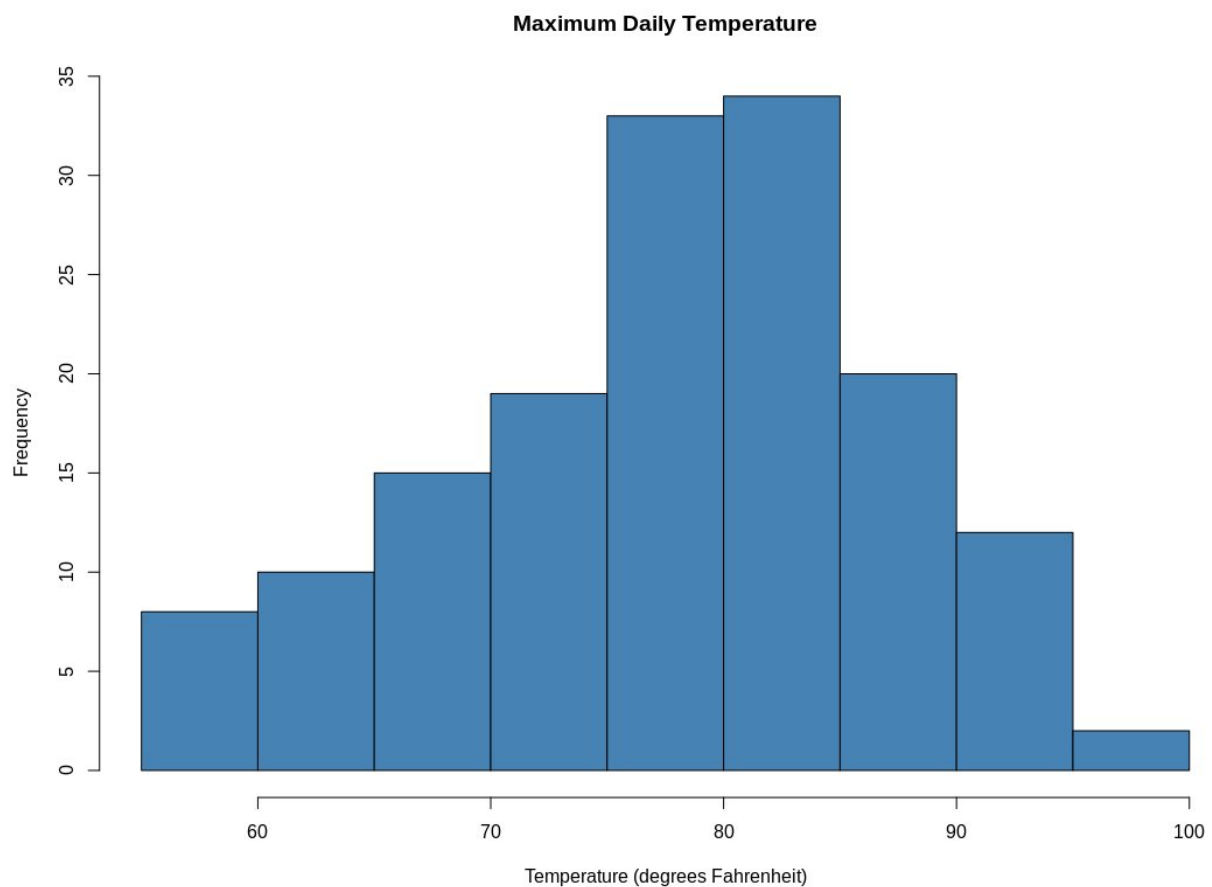
DATA VISUALIZATION USING HISTOGRAM PLOTTING FRAMEWORK

Histograms are standard in some academic fields, but they're usually reserved for the senior-most levels. These charts are best with highly precise or accurate numbers in R.

It ultimately provides a probability estimate of a variable — the period of time before a project's completion, for example. R provides a simple function for this as well.

histogram of frequency of ozone values in 'airquality' dataset

```
hist(airquality$Temp,col='steelblue',main='Maximum Daily Temperature',  
     xlab='Temperature (degrees Fahrenheit)')
```



VIVA QUESTIONS

1. When analyzing a histogram, what are some of the features to look for?

A histogram has an appearance similar to a vertical bar chart, but there are no gaps between the bars. Generally, a histogram will have bars of equal width.

2. What type of data is histogram usually used for?

The histogram is a popular graphing tool. It is used to summarize discrete or continuous data that are measured on an interval scale.

3. When will you use a histogram and when will you use a bar chart? Explain with an example.

The histogram is used to showcase a graphical presentation that represents the data in the form of frequency; whereas a bar chart is also a graphical representation of data and the information that is used for the comparison of two categories.

4. How do you make multiple plots to a single page layout in R?

To put multiple plots on the same graphics pages in R, you can use the graphics parameter `mfrow` or `mfc`.

5. How missing values and impossible values are represented in R?


In R, missing values are represented by the symbol `NA` (not available). Impossible values (domain errors like division by 0 et logs of negative numbers are represented by the symbol `NaN` (Not-A-Number). `NA` is used for both numeric and string data.

WEEK11

DATA VISUALIZATION USING LINE GRAPH PLOTTING, SCATTERPLOT PLOTTING FRAMEWORK

Plotting is a popular alternative to charting or graphing. It provides a unique visualization involving various dots. The most standard iteration — the scatter plot — tracks two continuous variables over the course of time. A basic application of the scatter plot involves tracking the height and weight of children throughout the years. Scatter plots are useful when trying to avoid misinformation in the visualization. Only use a plot if you're sure the audience is familiar with that type of chart, and always use it sparingly. When in doubt, go with one of your other options.

```
# Plot Ozone and Temperature measurements for only the month of September  
with(subset(airquality,Month==9),plot(Wind,Ozone,col='steelblue',pch=20,cex=1.5))  
title('Wind and Temperature in NYC in September of 1973')
```



VIVA QUESTIONS

1. What is a scatter plot?

A scatterplot is a type of data display that shows the relationship between two numerical variables.

2. What features might be visible in scatterplots?

Visually show the strength of the relationship between the variables.

3. Explain what should be done with suspected or missing data?

- Use deletion methods to eliminate missing data. The deletion methods only work for certain datasets where participants have missing fields.
- Use regression analysis to systematically eliminate data.
- Data scientists can use data imputation techniques.

4. For what type of data is scatter plot usually used for?

Scatter plots' primary uses are to observe and show relationships between two numeric variables. Relationships between variables can be described in many ways: positive or negative, strong or weak, linear or nonlinear. A scatter plot can also be useful for identifying other patterns in data.

WEEK12


APPLICATION TO ANALYZE STOCK MARKET USING R LANGUAGE.

The stock market of India fell by 41.2% to be precise during the period between January to April. At the starting of this year, most of the companies were at their peak and turned down to be 52 weeks low (considered as lowest ever in the year). Now, the stock market started recovering and the gain percentage in an average is 43% but still, there are pending recovery values that need to be retrieved. This analysis is taken sector-wise and is compared with the NSE (National Stock Exchange) indices movement. All these data and scrutiny (observation) is between Jan-01–2020 to Jun-09–2020. This analysis also reveals how the stocks and index performed during this period and which are the best and worst recovered companies.

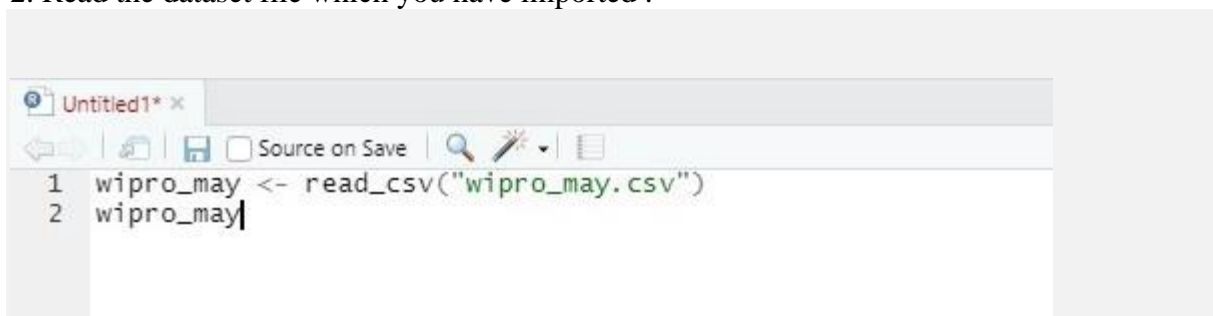
R as an Analysis Tool

R programming language is considered one of the powerful languages for Data Science and it is also very accessible for users. The software used for coding R is RStudio and there are also other platforms to code R which you can prefer. This language is widely used by Statisticians and Data miners for data analysis and for developing statistical software. Let's see how to do Stock Analysis with R :

1. Importing dataset using **Import Dataset** :



2. Read the dataset file which you have imported :



3. Coding for our required Analysis (I have extracted the data from NSE Website, the above one is just an example, now I'm coding with a different dataset)

This is an example with Wipro :

```

library(dplyr)
#Examining the lowest value of wipro
> min(wipro_index$Close)
162.35
#Examining the highest value of wipro
> max(wipro_index$Close)
257.2
#Calculating fall value of wipro
> max(wipro_index$Close) - min(wipro_index$Close)
94.85
#Calculating fall percentage of wipro
> 94.85 / max(wipro_index$Close) * 100
36.88
#Examining the recovery of wipro
> max(wipro_may$'Close Price')
226.45
#Calculating recovery value of wipro
> max(wipro_may$'Close Price') - min(wipro_index$Close)
64.1
#Calculating recovery percentage of wipro
> 64.1 / min(wipro_index$Close) * 100
39.48
#Calculating pending recovery value of wipro
> max(wipro_index$Close) - max(wipro_may$'Close Price')
30.75
#Calculating pending recovery percentage of wipro
> 30.75 / max(wipro_index$Close) * 100
11.96
#Summarizing our analysis of wipro
> wipro_index %>%
  summarize(high = max(wipro_index$Close) ,
    low = min(wipro_index$Close) ,
    fall_value = max(wipro_index$Close) - min(wipro_index$Close) ,
    fall_percentage = 94.85 / max(wipro_index$Close) * 100 ,
    recovery = max(wipro_may$`Close Price`) ,
    recovery_value = max(wipro_may$`Close Price`) - min(wipro_index$Close) ,
    recovery_percentage = 64.10 / min(wipro_index$Close) * 100 ,
    pending_recovery = max(wipro_index$Close) - max(wipro_may$`Close Price`) ,
    pending_percentage = 30.75 / max(wipro_index$Close) * 100 )
#Output
high low fall_value fall_percentage recovery recovery_value
257.2 162.35 94.85 36.88 226.45 64.1
recovery_percentage pending_recovery pending_percentage 39.4826 30.75
11.96

```

From this coding, we can generate our specified result which we want to form the dataset. After getting our data from R we can use Excel for a neat tabulation for our Analysis. In this program, I have used functions including “min,” ”max,” “summarize.”

min — “min” function returns the minimum value of a vector or column.

max — It returns the highest or the maximum value of a vector or column.

summarize — The “summarize” function reduces a data frame into a summary of just one vector or value.

Tabulation of our Data

| Stock Prices During COVID - 19 | | | | | | | | | | |
|--------------------------------|------------|-----------|----------|------------------------|------------------------|-----------|-----------------------|-------------------------|------------------------|--------|
| Companies | Symbol | High | Low | Fall Value(High - Low) | % (Value / High * 100) | Recovery | Value(Recovery - Low) | %(Recovery / Low * 100) | Pending Recovery Value | % |
| Nifty 50 | NIFTY50 | 12,362.30 | 7,610.25 | 4,752.05 | 38.44 | 10,167.25 | 2,557.20 | 33.60 | 2,195.05 | 21.59% |
| Wipro | WIPRO | 257.20 | 162.35 | 94.85 | 36.88 | 226.45 | 64.10 | 39.48 | 30.75 | 13.58% |
| Tata Motors | TATAMOTORS | 200.35 | 65.30 | 135.05 | 67.41 | 115.45 | 50.15 | 76.80 | 84.90 | 73.54% |
| SBI | SBIN | 339.30 | 175.50 | 163.80 | 48.28 | 187.80 | 12.30 | 7.01 | 151.50 | 80.67% |
| Sun Pharma | SUNPHARMA | 489.90 | 324.50 | 165.40 | 33.76 | 500.75 | 176.25 | 54.31 | 10.85 | 2.17% |
| Reliance | RELIANCE | 1,581.00 | 884.05 | 696.95 | 44.08 | 1,581.70 | 697.65 | 78.92 | 0.70 | 0.04% |
| L & T | LT | 1,370.20 | 707.90 | 662.30 | 48.34 | 961.35 | 253.45 | 35.80 | 408.85 | 42.53% |
| ITC | ITC | 243.25 | 147.25 | 96.00 | 39.47 | 200.55 | 53.30 | 36.20 | 42.70 | 21.29% |
| Airtel | BHARTIARTL | 565.00 | 404.05 | 160.95 | 28.49 | 598.80 | 194.75 | 48.20 | 33.80 | 5.64% |
| Gail | GAIL | 131.65 | 69.40 | 62.25 | 47.28 | 104.75 | 35.35 | 50.94 | 26.90 | 25.68% |
| Asian Paints | ASIANPAINT | 1,893.70 | 1,498.45 | 395.25 | 20.87 | 1,716.55 | 218.10 | 14.56 | 177.15 | 10.32% |

Graph comparison of Values


Data visualization is very important in the financial world, especially for Analysis. Instead of creating a tabulation, a simple graph or any way of representation is easy to understand and it is very useful for comparison too. In this blog, I’ve compared the values of companies using line graphs for easy understanding.

1. Nifty 50




From this graph, we can understand that Nifty 50 index was at its peak on 14-Jan-2020 and had a fall price of 7,610.250 on 23-Mar-2020. We can also observe that the price of 12,362.30 came down to 7,610.250 in just 68 days which is 38.44% from its peak value. The next thing we can observe from this graph is the recovery value of Nifty 50. The stock values of Nifty 50 zoomed from 7,610 to 10,167 in only 78 days and had a jump percentage of 33.6 %. The last thing about this graph is that even though the prices spiked in 78 days it couldn't make up to its peak price of 12,362. So, we can conclude that the pending recovery value of Nifty 50 is still 2,195 and needs to shine up for a percentage of 21.59%.

2. Wipro and ITC (Indian Tobacco Company)




This is a graph comparing stock prices of two companies, Wipro and ITC. We can observe that the movement of both companies is relatively similar during this period. The stock price of Wipro in January was 257.2, had a fall of 94.85, and reached a stock price of 162. The percentage of fall value is 36.88%. Soon it recovered in June and zoomed to a price of 226 and had a recovery rate of 39.4%. When we compare the peak price of Wipro with the recovery price, the pending recovery value is 30.75 and has to rise by 13.58%. Coming to ITC, its peak price was 243.25 and declined to a price of 147.25 during this period. This company had a fall value of 96 and declined by 39.47%. The stock prices are appreciated to 200.55 between March and June, the recovery rate was 36.20%. In case of pending recovery value, Wipro has a less pending recovery rate when compared to ITC which is at 21.29%. We can say that Wipro performed well during this period as it has less pending recovery rate when compared to ITC.

3. Tata Motors and Gail



This graph compares the stock prices of two companies Tata Motors and Gail. We can notice that the peak price of these two companies differs but in case of declined price, they are similar. The peak price of Tata Motors was 200.35 and the declined price is 65.30. From these prices, we can observe that Tata Motors has phenomenal fall value and rate. The fall value of Tata Motors is 135.05 and declined percentage is 64.41%. In Spite of having drastic depreciation in prices Tata Motors recovered in a very short span of time. It zoomed to a price of 115.45 from a declined price of 65.30. The recovery rate is 76.80%. But still, the pending recovery value is 84.90 with a rate of 73.54%. Coming to Gail the price movements are not so phenomenal when compared to Tata Motors. The peak price was 131.65 and the declined price is 69.40 which is more similar to Tata Motors. The fall value and rate of Gail are 62.25 and 42.28% respectively. The recovery price is 104.75 which can be considered as a good comeback when compared to Tata Motors. The pending recovery value of Gail is 26.90 with a percentage of 25.68%. This graph comparison clearly reveals that the price movements of Tata Motors during this period were very volatile.

4. Reliance, L & T and Asian Paints



This graph comprises the stock prices of three companies namely Reliance, L & T, and Asian Paints. The stock price of Reliance was its peak value of 1,581 and had a price depreciation to 884.05. But the recovery rate was phenomenally high which is higher than its peak price. The recovery price is 1,581.70, an increased rate of 0.7% from its peak price. Obviously there is no pending recovery rate as the price exceeds its peak value. The highest stock price of L & T was 1,370.20. The fall price and value of L & T are 707.90 and 662.30 respectively. The down rate of L & T is 48.34%. The recovery price of L & T is 961.35 and the rate is 35.80. The pending recovery value of L & T is 408.85. From this we can say that the performance of L & T is poor when compared to L & T. Coming to Asian Paints, considering all the peak prices Asian Paints is top among that. The peak price of Asian Paints was 1,893.70 and the declined price is 1,498.45. The recovery price is 1,716.55, from this we can note that the pending value and rate of Asian Paints would be lower than L & T. The pending recovery value and rate of Asian Paints are 177.15 and 10.32%. Like this, we can predict the values easily in charts or any type of representation without doing any calculations.

VIVA QUESTIONS

1. Is R good for Stock analysis?

Yes, R has excellent packages for analyzing stock data, so I feel there should be a “translation” of the post for using R for stock data analysis.

2. How do I get Stock data in R?

By using quantmod package. You can install it by typing the command “install.packages(“quantmod”)” in your R console.

3. How do you analyze data from the stock market?

The quantmod package for R is designed to assist the quantitative trader in the development, testing, and deployment of statistically based trading models.

4. Which app is best for stock market analysis?

MoneyControl, Stock Edge, Economic Times (ET) Markets

5. List out some of the functions that R provides?

Apline, abs, addmargins, aggregare and etc..

ADDITIONAL PROGRAMS

IMPLEMENT BASIC PROGRAMS IN R BY USING DATA STRUCTURES

```
> myString <- "Hello, World!"
> print ( myString)
[1] "Hello, World!"
```

Vectors:

```
apple <- c('red','green',"yellow")
print(apple)
print(class(apple))
[1] "red"  "green" "yellow"
[1] "character"
```

Lists:

```
list1 <- list(c(2,5,3),21.3,sin)
print(list1)
[[1]]
[1] 2 5 3
[[2]]
[1] 21.3
[[3]]
function (x) .Primitive("sin")
```

Matrices:

```
M=matrix(c('a','a','b','c','b','a'),
  nrow=2,ncol=3,byrow=TRUE)
print(M)
[,1] [,2] [,3]
[1,] "a"  "a"  "b"
[2,] "c"  "b"  "a"
```

Arrays:

```
a <- array(c('green','yellow'),dim = c(3,3,1))
print(a)
, , 1
[,1] [,2] [,3]
[1,] "green" "yellow" "green"
[2,] "yellow" "green" "yellow"
[3,] "green" "yellow" "green"
```

Data Frames:

```
BMI <- data.frame(
  gender = c("Male", "Male","Female"),
  height = c(152, 171.5, 165),
  weight = c(81,93, 78),
  Age = c(42,38,26)
)
print(BMI)
  gender height weight Age
1  Male  152.0    81  42
2  Male  171.5    93  38
3 Female  165.0    78  26
```

1. Write a R program to take input from the user (name and age) and display the values. Also print the version of R installation.

```
name = readline(prompt="Input your name: ")
age = readline(prompt="Input your age: ")
print(paste("My name is",name, "and I am",age ,"years old."))
print(R.version.string)
```

Output

```
Input your name: abc
Input your age: 25
[1] "My name is abc and I am 25 years old."
[1] "R version 3.4.4 (2018-03-15)"
```

2. Write a R program to get the details of the objects in memory.

```
name = "Python";
n1 = 10;
n2 = 0.5
nums = c(10, 20, 30, 40, 50, 60)
print(ls())
print("Details of the objects in memory:")
print(ls.str())
```

Output

```
[1] "n1" "n2" "name" "nums"
[1] "Details of the objects in memory:"
n1 : num 10
n2 : num 0.5
name : chr "Python"
nums : num [1:6] 10 20 30 40 50 60
```

3. Write a R program to create a sequence of numbers from 20 to 50 and find the mean of numbers from 20 to 60 and sum of numbers from 51 to 91.

```
print("Sequence of numbers from 20 to 50:")
print(seq(20,50))
print("Mean of numbers from 20 to 60:")
print(mean(20:60))
print("Sum of numbers from 51 to 91:")
print(sum(51:91))
```

Output

```
[1] "Sequence of numbers from 20 to 50:"
[1] 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44
[26] 45 46 47 48 49 50
[1] "Mean of numbers from 20 to 60:"
[1] 40
[1] "Sum of numbers from 51 to 91:"
```


[1] 2911

4. Write a R program to create a vector which contains 10 random integer values between -50 and +50.

```
v = sample(-50:50, 10, replace=TRUE)
print("Content of the vector:")
print("10 random integer values between -50 and +50:")
print(v)
```

Output

```
[1] "Content of the vector:"
[1] "10 random integer values between -50 and +50:"
[1] 31 -13 -21 42 49 -39 20 12 39 -2
```

5. Write a R program to get the first 10 Fibonacci numbers.

```
Fibonacci <- numeric(10)
Fibonacci[1] <- Fibonacci[2] <- 1
for (i in 3:10) Fibonacci[i] <- Fibonacci[i - 2] + Fibonacci[i - 1]
print("First 10 Fibonacci numbers:")
print(Fibonacci)
```

Output

```
[1] "First 10 Fibonacci numbers:"
[1] 1 1 2 3 5 8 13 21 34 55
```

6. Write a R program to get all prime numbers up to a given number (based on the sieve of Eratosthenes).

```
prime_numbers <- function(n) {
  if (n >= 2) {
    x = seq(2, n)
    prime_nums = c()
    for (i in seq(2, n)) {
      if (any(x == i)) {
        prime_nums = c(prime_nums, i)
        x = x[(x %% i) != 0], i
      }
    }
    return(prime_nums)
  }
  else
  {
    stop("Input number should be at least 2.")
  }
}
prime_numbers(12)
```

Output

```
[1] 2 3 5 7 11
```

7. Write a R program to print the numbers from 1 to 100 and print "Fizz" for multiples of 3, print "Buzz" for multiples of 5, and print "FizzBuzz" for multiples of both

```
for (n in 1:100) {  
  if (n %% 3 == 0 & n %% 5 == 0) {print("FizzBuzz")}  
  else if (n %% 3 == 0) {print("Fizz")}  
  else if (n %% 5 == 0) {print("Buzz")}  
  else print(n)  
}
```

Output

```
[1] 1  
[1] 2  
[1] "Fizz"  
[1] 4  
[1] "Buzz"  
[1] "Fizz"  
[1] 7  
[1] 8  
[1] "Fizz"  
[1] "Buzz"  
[1] 11  
[1] "Fizz"  
[1] 13  
[1] 14  
[1] "FizzBuzz"  
[1] 16  
[1] 17  
[1] "Fizz"  
[1] 19  
[1] "Buzz"  
[1] "Fizz"  
[1] 22  
[1] 23  
[1] "Fizz"  
[1] "Buzz"  
[1] 26  
[1] "Fizz"  
[1] 28  
[1] 29  
[1] "FizzBuzz"  
[1] 31  
[1] 32  
[1] "Fizz"  
[1] 34  
[1] "Buzz"
```

```
[1] "Fizz"
[1] 37
[1] 38
[1] "Fizz"
[1] "Buzz"
[1] 41
[1] "Fizz"
[1] 43
[1] 44
[1] "FizzBuzz"
[1] 46
[1] 47
[1] "Fizz"
[1] 49
[1] "Buzz"
[1] "Fizz"
[1] 52
[1] 53
[1] "Fizz"
[1] "Buzz"
[1] 56
[1] "Fizz"
[1] 58
[1] 59
[1] "FizzBuzz"
[1] 61
[1] 62
[1] "Fizz"
[1] 64
[1] "Buzz"
[1] "Fizz"
[1] 67
[1] 68
[1] "Fizz"
[1] "Buzz"
[1] 71
[1] "Fizz"
[1] 73
[1] 74
[1] "FizzBuzz"
[1] 76
[1] 77
[1] "Fizz"
[1] 79
[1] "Buzz"
[1] "Fizz"
[1] 82
[1] 83
[1] "Fizz"
[1] "Buzz"
```

```
[1] 86
[1] "Fizz"
[1] 88
[1] 89
[1] "FizzBuzz"
[1] 91
[1] 92
[1] "Fizz"
[1] 94
[1] "Buzz"
[1] "Fizz"
[1] 97
[1] 98
[1] "Fizz"
[1] "Buzz"
```

8. Write a R program to extract first 10 english letter in lower case and last 10 letters in upper case and extract letters between 22nd to 24th letters in upper case.

```
print("First 10 letters in lower case:")
t = head(letters, 10)
print(t)
print("Last 10 letters in upper case:")
t = tail(LETTERS, 10)
print(t)
print("Letters between 22nd to 24th letters in upper case:")
e = tail(LETTERS[22:24])
print(e)
```

Output

```
[1] "First 10 letters in lower case:"
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j"
[1] "Last 10 letters in upper case:"
[1] "Q" "R" "S" "T" "U" "V" "W" "X" "Y" "Z"
[1] "Letters between 22nd to 24th letters in upper case:"
[1] "V" "W" "X"
```

9. Write a R program to find the factors of a given number.

```
print_factors = function(n) {
  print(paste("The factors of",n,"are:"))
  for(i in 1:n) {
    if((n %% i) == 0) {
      print(i)
    }
  }
}
print_factors(4)
```

```
print_factors(7)
print_factors(12)
```

Output

```
[1] "The factors of 4 are:"
[1] 1
[1] 2
[1] 4
[1] "The factors of 7 are:"
[1] 1
[1] 7
[1] "The factors of 12 are:"
[1] 1
[1] 2
[1] 3
[1] 4
[1] 6
[1] 12
```

10. Write a R program to find the maximum and the minimum value of a given vector.

```
nums = c(10, 20, 30, 40, 50, 60)
print('Original vector:')
print(nums)
print(paste("Maximum value of the said vector:",max(nums)))
print(paste("Minimum value of the said vector:",min(nums)))
```

Output

```
[1] "Original vector:"
[1] 10 20 30 40 50 60
[1] "Maximum value of the said vector: 60"
[1] "Minimum value of the said vector: 10"
```

11. Write a R program to get the unique elements of a given string and unique numbers of vector

str1 = "The quick brown fox jumps over the lazy dog."

```
print("Original vector(string)")
print(str1)
print("Unique elements of the said vector:")
print(unique(tolower(str1)))
nums = c(1, 2, 2, 3, 4, 4, 5, 6)
print("Original vector(number)")
print(nums)
```

```
print("Unique elements of the said vector:")
print(unique(nums))
```

Output

```
[1] "Original vector(string)"
[1] "The quick brown fox jumps over the lazy dog."
[1] "Unique elements of the said vector:"
[1] "the quick brown fox jumps over the lazy dog."
[1] "Original vector(number)"
[1] 1 2 2 3 4 4 5 6
[1] "Unique elements of the said vector:"
[1] 1 2 3 4 5 6
```

12. Write a R program to create three vectors a,b,c with 3 integers. Combine the three vectors to become a 3×3 matrix where each column represents a vector. Print the content of the matrix

```
a<-c(1,2,3)
b<-c(4,5,6)
c<-c(7,8,9)
m<-cbind(a,b,c)
print("Content of the said matrix:")
print(m)
```

Output

```
[1] "Content of the said matrix:"
      a b c
[1,] 1 4 7
[2,] 2 5 8
[3,] 3 6 9
```

SUMMARIZING THE STATISTICS

1. summary(iris)

Output:

| Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species |
|---------------|---------------|---------------|---------------|---------------|
| Min. :4.300 | Min. :2.000 | Min. :1.000 | Min. :0.100 | setosa :50 |
| 1st Qu.:5.100 | 1st Qu.:2.800 | 1st Qu.:1.600 | 1st Qu.:0.300 | versicolor:50 |
| Median :5.800 | Median :3.000 | Median :4.350 | Median :1.300 | virginica :50 |
| Mean :5.843 | Mean :3.057 | Mean :3.758 | Mean :1.199 | |
| 3rd Qu.:6.400 | 3rd Qu.:3.300 | 3rd Qu.:5.100 | 3rd Qu.:1.800 | |

Max. :7.900 Max. :4.400 Max. :6.900 Max. :2.500

2. **str(mtcars)**

Output:

'data.frame': 32 obs. of 11 variables:

\$ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...

\$ cyl : num 6 6 4 6 8 6 8 4 4 6 ...

\$ disp: num 160 160 108 258 360 ...

\$ hp : num 110 110 93 110 175 105 245 62 95 123 ...

\$ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...

\$ wt : num 2.62 2.88 2.32 3.21 3.44 ...

\$ qsec: num 16.5 17 18.6 19.4 17 ...

\$ vs : num 0 0 1 1 0 1 0 1 1 1 ...

\$ am : num 1 1 1 0 0 0 0 0 0 0 ...

\$ gear: num 4 4 4 3 3 3 3 4 4 4 ...

\$ carb: num 4 4 1 1 2 1 4 2 2 4 ...

3. **head(mtcars)**

Output:

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| Valiant | 18.1 | 6 | 225 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |

4. **tail(mtcars)**

Output:

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|---------------|------|-----|-------|----|------|-------|------|----|----|------|------|
| Porsche 914-2 | 26.0 | 4 | 120.3 | 91 | 4.43 | 2.140 | 16.7 | 0 | 1 | 5 | 2 |

| | | | | | | | | | | | |
|----------------|------|---|-------|-----|------|-------|------|---|---|---|---|
| Lotus Europa | 30.4 | 4 | 95.1 | 113 | 3.77 | 1.513 | 16.9 | 1 | 1 | 5 | 2 |
| Ford Pantera L | 15.8 | 8 | 351.0 | 264 | 4.22 | 3.170 | 14.5 | 0 | 1 | 5 | 4 |
| Ferrari Dino | 19.7 | 6 | 145.0 | 175 | 3.62 | 2.770 | 15.5 | 0 | 1 | 5 | 6 |
| Maserati Bora | 15.0 | 8 | 301.0 | 335 | 3.54 | 3.570 | 14.6 | 0 | 1 | 5 | 8 |
| Volvo 142E | 21.4 | 4 | 121.0 | 109 | 4.11 | 2.780 | 18.6 | 1 | 1 | 4 | 2 |

5. **names(mtcars)**

Output:

```
[1] "mpg" "cyl" "disp" "hp" "drat" "wt" "qsec" "vs" "am" "gear" "carb"
```

6. **nrow(mtcars)**

Output:

```
[1] 32
```

7. **aggregate(Sepal.Length~Species,iris,mean)**

Output:

| | Species | Sepal.Length |
|---|------------|--------------|
| 1 | setosa | 5.006 |
| 2 | versicolor | 5.936 |
| 3 | virginica | 6.588 |

8. **fix(iris)**

Output:

| | Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species | var6 |
|----|--------------|-------------|--------------|-------------|---------|------|
| 1 | 5.1 | 3.5 | 1.4 | 0.2 | setosa | |
| 2 | 4.9 | 3 | 1.4 | 0.2 | setosa | |
| 3 | 4.7 | 3.2 | 1.3 | 0.2 | setosa | |
| 4 | 4.6 | 3.1 | 1.5 | 0.2 | setosa | |
| 5 | 5 | 3.6 | 1.4 | 0.2 | setosa | |
| 6 | 5.4 | 3.9 | 1.7 | 0.4 | setosa | |
| 7 | 4.6 | 3.4 | 1.4 | 0.3 | setosa | |
| 8 | 5 | 3.4 | 1.5 | 0.2 | setosa | |
| 9 | 4.4 | 2.9 | 1.4 | 0.2 | setosa | |
| 10 | 4.9 | 3.1 | 1.5 | 0.1 | setosa | |
| 11 | 5.4 | 3.7 | 1.5 | 0.2 | setosa | |
| 12 | 4.8 | 3.4 | 1.6 | 0.2 | setosa | |
| 13 | 4.8 | 3 | 1.4 | 0.1 | setosa | |
| 14 | 4.3 | 3 | 1.1 | 0.1 | setosa | |
| 15 | 5.8 | 4 | 1.2 | 0.2 | setosa | |
| 16 | 5.7 | 4.4 | 1.5 | 0.4 | setosa | |
| 17 | 5.4 | 3.9 | 1.3 | 0.4 | setosa | |
| 18 | 5.1 | 3.5 | 1.4 | 0.3 | setosa | |
| 19 | 5.7 | 3.8 | 1.7 | 0.3 | setosa | |

9. `sepalsub<- subset(iris,Sepal.Length>7)`

sepalsub

Output:

| | Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species | var6 |
|-----------|--------------|-------------|--------------|-------------|-----------|------|
| 103 A> | 7.1 | 3.0 | 5.9 | 2.1 | virginica | <N |
| 106 A> | 7.6 | 3.0 | 6.6 | 2.1 | virginica | <N |
| 108 A> | 7.3 | 2.9 | 6.3 | 1.8 | virginica | <N |
| 110 A> | 7.2 | 3.6 | 6.1 | 2.5 | virginica | <N |
| 118 A> | 7.7 | 3.8 | 6.7 | 2.2 | virginica | <N |
| 119 A> | 7.7 | 2.6 | 6.9 | 2.3 | virginica | <N |
| 123 A> | 7.7 | 2.8 | 6.7 | 2.0 | virginica | <N |
| 126 A> | 7.2 | 3.2 | 6.0 | 1.8 | virginica | <N |

| | | | | | | |
|-----------|-----|-----|-----|-----|-----------|----|
| 130 A> | 7.2 | 3.0 | 5.8 | 1.6 | virginica | <N |
| 131 A> | 7.4 | 2.8 | 6.1 | 1.9 | virginica | <N |
| 132 A> | 7.9 | 3.8 | 6.4 | 2.0 | virginica | <N |
| 136 A> | 7.7 | 3.0 | 6.1 | 2.3 | virginica | <N |

BINOMIAL DISTRIBUTION

1. **choose(10,3)*((1/6)^3*(5/6)^7)**

Output:

[1] 0.1550454

2. **dbinom(3,size=10,prob=(1/6))**

Output:

[1] 0.1550454

3. **choose(10,0)*((1/6)^0*(5/6)^10) +**

choose(10,1)*((1/6)^1*(5/6)^9)+

choose(10,2)*((1/6)^2*(5/6)^8)+

choose(10,3)*((1/6)^3*(5/6)^7)

Output:

[1] 0.9302722

4. **pbinom(3,size=10,prob=(1/6),lower=T)**

Output:

[1] 0.9302722

5. **pbinom(3,size=10,prob=(1/6),lower=F)**

Output:

[1] 0.06972784

6. **dbinom(4,size=12,prob=0.2)**

Output:

[1] 0.1328756

7. **dbinom(0,size=12,prob=0.2)+**

dbinom(1,size=12,prob=0.2)+

dbinom(2,size=12,prob=0.2)+

dbinom(3,size=12,prob=0.2)+

dbinom(4,size=12,prob=0.2)

Output:

[1] 0.9274445

8. **pbinom(4,size=12,prob=0.2)**

Output:

[1] 0.9274445

9. **x <- pbinom(26,51,0.5)**

print(x)

Output:


[1] 0.610116

10. **x <- seq(0,50,by = 1)**

y <- dbinom(x,50,0.5)

plot(x,y)

Output:



```
11. x <- qbinom(0.25,51,1/2)
```

```
print(x)
```

Output:

```
[1] 23
```

```
12. x <- rbinom(8,150,.4)
```

```
print(x)
```

Output:

```
[1] 60 71 57 60 62 62 50 59
```

POISSON DISTRIBUTION

```
1. ppois(16,lambda = 12)
```

Output:

```
[1] 0.898709
```

```
2. ppois(16,lambda = 12,lower=F)
```

Output:

```
[1] 0.101291
```

```
3. rpois(16,lambda = 12)
```

Output:

```
[1] 12 8 12 10 13 8 11 13 12 11 15 12 10 11 14 11
```

```
4. dpois(16,lambda = 12)
```

Output:

```
[1] 0.05429334
```

NORMAL DISTRIBUTION

1. `pnorm(84,mean=72,sd=15.2,lower.tail=FALSE)`

Output:


[1] 0.2149176

2. `x <- seq(-10, 10, by = .1)`

`y <- dnorm(x, mean = 2.5, sd = 0.5)`

`plot(x,y)`

Output:




3. `x <- seq(-10,10,by = .2)`

`y <- pnorm(x, mean = 2.5, sd = 2)`

`plot(x,y)`

Output:




```
4. x <- seq(0, 1, by = 0.02)
```

```
   y <- qnorm(x, mean = 2, sd = 1)
```

```
   plot(x,y)
```


Output:



```
5. y <- rnorm(50)
```

```
   hist(y, main = "Normal DIstribution")
```

Output:



LINEAR REGRESSION

Write a program to implement linear and logistic regression


```
head(cars)
```

Output:

| | speed | dist |
|---|-------|------|
| 1 | 4 | 2 |
| 2 | 4 | 10 |
| 3 | 7 | 4 |
| 4 | 7 | 22 |
| 5 | 8 | 16 |
| 6 | 9 | 10 |

```
scatter.smooth(x=cars$speed,y=cars$dist,main="dist~speed")
```

Output:



```
linearMod<-lm(dist~speed,data=cars)
```

```
print(linearMod)
```

Output:

Coefficients:

| | |
|-------------|-------|
| (Intercept) | speed |
| -17.579 | 3.932 |

summary(linearMod)

Output:

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|--------|--------|-------|--------|
| -29.069 | -9.525 | -2.272 | 9.215 | 43.201 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|--------------|
| (Intercept) | -17.5791 | 6.7584 | -2.601 | 0.0123 * |
| speed | 3.9324 | 0.4155 | 9.464 | 1.49e-12 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 15.38 on 48 degrees of freedom

Multiple R-squared: 0.6511, Adjusted R-squared: 0.6438

F-statistic: 89.57 on 1 and 48 DF, p-value: 1.49e-12

LOGISTIC REGRESSION

```
input<-mtcars[,c("am","cyl","hp","wt")]
```

```
print(input)
```

Output:

| | am | cyl | hp | wt |
|-------------------|----|-----|-----|-------|
| Mazda RX4 | 1 | 6 | 110 | 2.620 |
| Mazda RX4 Wag | 1 | 6 | 110 | 2.875 |
| Datsun 710 | 1 | 4 | 93 | 2.320 |
| Hornet 4 Drive | 0 | 6 | 110 | 3.215 |
| Hornet Sportabout | 0 | 8 | 175 | 3.440 |
| Valiant | 0 | 6 | 105 | 3.460 |
| Duster 360 | 0 | 8 | 245 | 3.570 |
| Merc 240D | 0 | 4 | 62 | 3.190 |
| Merc 230 | 0 | 4 | 95 | 3.150 |
| Merc 280 | 0 | 6 | 123 | 3.440 |

| | |
|------------|---------------|
| Merc 280C | 0 6 123 3.440 |
| Merc 450SE | 0 8 180 4.070 |
| Merc 450SL | 0 8 180 3.730 |

```
input<-mtcars[,c("am","cyl","hp","wt")]
am.data=glm(formula=am~cyl+hp+wt,data=input,family=binomial)
print(summary(am.data))
```

Output:

Deviance Residuals:

| Min | 1Q | Median | 3Q | Max |
|----------|----------|----------|---------|---------|
| -2.17272 | -0.14907 | -0.01464 | 0.14116 | 1.27641 |

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|-------------|----------|------------|---------|----------|
| (Intercept) | 19.70288 | 8.11637 | 2.428 | 0.0152 * |
| cyl | 0.48760 | 1.07162 | 0.455 | 0.6491 |
| hp | 0.03259 | 0.01886 | 1.728 | 0.0840 . |
| wt | -9.14947 | 4.15332 | -2.203 | 0.0276 * |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 43.2297 on 31 degrees of freedom

Residual deviance: 9.8415 on 28 degrees of freedom

AIC: 17.841

Number of Fisher Scoring iterations: 8